

Chapter 10

Pointers

A *pointer* is a variable that represents the *location* (rather than the *value*) of a data item, such as a variable or an array element. Pointers are used frequently in C, as they have a number of useful applications. For example, pointers can be used to pass information back and forth between a function and its reference point. In particular, pointers provide a way to return multiple data items from a function via function arguments. Pointers also permit references to other functions to be specified as arguments to a given function. This has the effect of passing functions as arguments to the given function.

Pointers are also closely associated with arrays and therefore provide an alternate way to access individual array elements. Moreover, pointers provide a convenient way to represent multidimensional arrays, allowing a single multidimensional array to be replaced by a lower-dimensional array of pointers. This feature permits a group of strings to be represented within a single array, though the individual strings may differ in length.

10.1 FUNDAMENTALS

Within the computer's memory, every stored data item occupies one or more contiguous memory cells (i.e., adjacent words or bytes). The number of memory cells required to store a data item depends on the type of data item. For example, a single character will typically be stored in one byte (8 bits) of memory; an integer usually requires two contiguous bytes; a floating-point number may require four contiguous bytes; and a double-precision quantity may require eight contiguous bytes. (See Chap. 2 and Appendix D.)

Suppose *v* is a variable that represents some particular data item. The compiler will automatically assign memory cells for this data item. The data item can then be accessed if we know the location (i.e., the *address*) of the first memory cell.* The address of *v*'s memory location can be determined by the expression *&v*, where *&* is a unary operator, called the *address operator*, that evaluates the address of its operand.

Now let us assign the address of *v* to another variable, *pv*. Thus,

```
pv = &v
```

This new variable is called a *pointer* to *v*, since it "points" to the location where *v* is stored in memory. Remember, however, that *pv* represents *v*'s *address*, not its *value*. Thus, *pv* is referred to as a *pointer variable*. The relationship between *pv* and *v* is illustrated in Fig. 10.1.



Fig. 10.1 Relationship between *pv* and *v* (where *pv* = *&v* and *v* = **pv*)

The data item represented by *v* (i.e., the data item stored in *v*'s memory cells) can be accessed by the expression **pv*, where *** is a unary operator, called the *indirection operator*, that operates only on a pointer

* Adjacent memory cells within a computer are numbered consecutively, from the beginning to the end of the memory area. The number associated with each memory cell is known as the memory cell's *address*. Most computers use a hexadecimal numbering system to designate the addresses of consecutive memory cells, though some computers use an octal numbering system (see Appendix A).

variable. Therefore, `*pv` and `v` both represent the same data item (i.e., the contents of the same memory cells). Furthermore, if we write `pv = &v` and `u = *pv`, then `u` and `v` will both represent the same value; i.e., the value of `v` will indirectly be assigned to `u`. (It is assumed that `u` and `v` are of the same data type.)

EXAMPLE 10.1 Shown below is a simple program that illustrates the relationship between two integer variables, their corresponding addresses and their associated pointers.

```
#include <stdio.h>

main()
{
    int u = 3;
    int v;
    int *pu;      /* pointer to an integer */
    int *pv;      /* pointer to an integer */

    pu = &u;      /* assign address of u to pu */
    v = *pu;      /* assign value of u to v */
    pv = &v;      /* assign address of v to pv */

    printf("\nu=%d  &u=%X  pu=%X  *pu=%d", u, &u, pu, *pu);
    printf("\n\nv=%d  &v=%X  pv=%X  *pv=%d", v, &v, pv, *pv);
}
```

Note that `pu` is a pointer to `u`, and `pv` is a pointer to `v`. Therefore `pu` represents the address of `u`, and `pv` represents the address of `v`. (Pointer declarations will be discussed in the next section.)

Execution of this program results in the following output.

```
u=3  &u=F8E  pu=F8E  *pu=3
v=3  &v=F8C  pv=F8C  *pv=3
```

In the first line, we see that `u` represents the value 3, as specified in the declaration statement. The address of `u` is determined automatically by the compiler as F8E (hexadecimal). The pointer `pu` is assigned this value; hence, `pu` also represents the (hexadecimal) address F8E. Finally, the value to which `pu` points (i.e., the value stored in the memory cell whose address is F8E) is 3, as expected.

Similarly, the second line shows that `v` also represents the value 3. This is expected, since we have assigned the value `*pu` to `v`. The address of `v`, and hence the value of `pv`, is F8C. Notice that `u` and `v` have different addresses. And finally, we see that the value to which `pv` points is 3, as expected.

The relationships between `pu` and `u`, and `pv` and `v`, are shown in Fig. 10.2. Note that the memory locations of the pointer variables (i.e., address EC7 for `pu`, and EC5 for `pv`) are not displayed by the program.

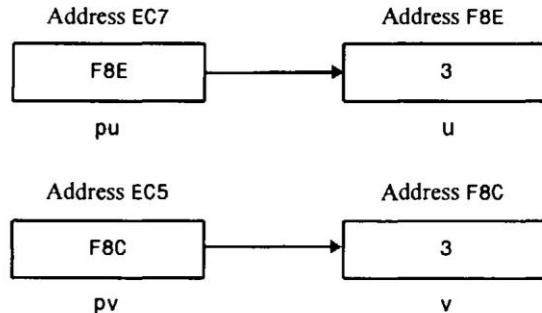


Fig. 10.2

The unary operators `&` and `*` are members of the same precedence group as the other unary operators, i.e., `-`, `++`, `--`, `!`, `sizeof` and `(type)`, which were presented in Chap. 3. Remember that this group of operators has a higher precedence than the groups containing the arithmetic operators, and that the associativity of the unary operators is right to left (see Appendix C).

The address operator (`&`) must act upon operands that are associated with unique addresses, such as ordinary variables or single array elements. Thus *the address operator cannot act upon arithmetic expressions*, such as `2 * (u + v)`.

The indirection operator (`*`) can only act upon operands that are pointers (e.g., pointer variables). However, if `pv` points to `v` (i.e., `pv = &v`), then an expression such as `*pv` can be used interchangeably with its corresponding variable `v`. Thus, an indirect reference (e.g., `*pv`) can appear in place of an ordinary variable (e.g., `v`) within a more complicated expression.

EXAMPLE 10.2 Consider the simple C program shown below.

```
#include <stdio.h>

main()
{
    int u1, u2;
    int v = 3;
    int *pv;           /* pv points to v */
    u1 = 2 * (v + 5);      /* ordinary expression */
    pv = &v;
    u2 = 2 * (*pv + 5);    /* equivalent expression */
    printf("\nu1=%d    u2=%d", u1, u2);
}
```

This program involves the use of two integer expressions. The first, `2 * (v + 5)`, is an ordinary arithmetic expression whereas the second, `2 * (*pv + 5)`, involves the use of a pointer. The expressions are equivalent, since `v` and `*pv` each represent the same integer value.

The following output is generated when the program is executed.

```
u1=16    u2=16
```

An indirect reference can also appear on the left side of an assignment statement. This provides another method for assigning a value to a variable or an array element.

EXAMPLE 10.3 A simple C program is shown below.

```
#include <stdio.h>

main()
{
    int v = 3;
    int *pv;

    pv = &v;           /* pv points to v */
    printf("\n*pv=%d    v=%d", *pv, v);

    *pv = 0;           /* reset v indirectly */
    printf("\n\n*pv=%d    v=%d", *pv, v);
}
```

The program begins by assigning an initial value of 3 to the integer variable *v*, and then assigns the address of *v* to the pointer variable *pv*. Thus, *pv* becomes a pointer to *v*. The expression **pv* therefore represents the value 3. The first *printf* statement is intended to illustrate this by displaying the current values of **pv* and *v*.

Following the first *printf* statement, the value of **pv* is reset to 0. Therefore, *v* will be reassigned the value 0. This is illustrated by the second *printf* statement, which causes the new values of **pv* and *v* to be displayed.

When the program is executed, the following output is generated.

```
*pv=3    v=3  
*pv=0    v=0
```

Thus, the value of *v* has been altered by assigning a new value to **pv*.

Pointer variables can point to numeric or character variables, arrays, functions or other pointer variables. (They can also point to certain other data structures that will be discussed later in this book.) Thus, a pointer variable can be assigned the address of an ordinary variable (e.g., *pv* = *&v*). Also, a pointer variable can be assigned the value of another pointer variable (e.g., *pv* = *px*), provided both pointer variables point to data items of the same type. Moreover, a pointer variable can be assigned a *null* (zero) value, as explained in Sec. 10.2 below. On the other hand, *ordinary* variables *cannot* be assigned arbitrary addresses (i.e., an expression such as *&x* cannot appear on the left-hand side of an assignment statement).

Section 10.5 presents additional information concerning those operations that can be carried out on pointers.

10.2 POINTER DECLARATIONS

Pointer variables, like all other variables, must be declared before they may be used in a C program. The interpretation of a pointer declaration differs, however, from the interpretation of other variable declarations. When a pointer variable is declared, the variable name must be preceded by an asterisk (*). This identifies the fact that the variable is a pointer. The data type that appears in the declaration refers to the *object* of the pointer, i.e., the data item that is stored in the address represented by the pointer, rather than the pointer itself.

Thus, a pointer declaration may be written in general terms as

```
data-type *ptvar;
```

where *ptvar* is the name of the pointer variable, and *data-type* refers to the data type of the pointer's object. Remember that an asterisk must precede *ptvar*.

EXAMPLE 10.4 A C program contains the following declarations.

```
float u, v;  
float *pv;
```

The first line declares *u* and *v* to be floating-point variables. The second line declares *pv* to be a pointer variable whose object is a floating-point quantity; i.e., *pv* points to a floating-point quantity. Note that *pv* represents an *address*, not a floating-point quantity. (Some additional pointer declarations are shown in Examples 10.1 to 10.3.)

Within a variable declaration, a pointer variable can be initialized by assigning it the address of another variable. Remember that the variable whose address is assigned to the pointer variable must have been declared earlier in the program.

EXAMPLE 10.5 A C program contains the following declarations.

```
float u, v;  
float *pv = &v;
```

The variables *u* and *v* are declared to be floating-point variables and *pv* is declared as a pointer variable that points to a floating-point quantity, as in Example 10.4. In addition, the address of *v* is initially assigned to *pv*.

This terminology can be confusing. Remember that these declarations are equivalent to writing

```
float u, v;      /* floating-point variable declarations */
float *pv;        /* pointer variable declaration */
. . .
pv = &v;          /* assign v's address to pv */
```

Note that an asterisk is not included in the assignment statement.

In general, it does not make sense to assign an integer value to a pointer variable. An exception, however, is an assignment of 0, which is sometimes used to indicate some special condition. In such situations the recommended programming practice is to define a symbolic constant **NULL** which represents 0, and to use **NULL** in the pointer initialization. This practice emphasizes the fact that the zero assignment represents a special situation.

EXAMPLE 10.6 A C program contains the following symbolic constant definitions and array declarations.

```
#define NULL 0
float u, v;
float *pv = NULL;
```

The variables *u* and *v* are declared to be floating-point variables and *pv* is declared as a pointer variable that points to a floating-point quantity. In addition, *pv* is initially assigned a value of 0 to indicate some special condition dictated by the logic of the program (which is not shown in this example). The use of the symbolic constant **NULL** suggests that this initial assignment is something other than the assignment of an ordinary integer value.

We will see other kinds of pointer declarations later in this chapter.

10.3 PASSING POINTERS TO A FUNCTION

Pointers are often passed to a function as arguments. This allows data items within the calling portion of the program to be accessed by the function, altered within the function, and then returned to the calling portion of the program in altered form. We refer to this use of pointers as passing arguments by *reference* (or by *address* or by *location*), in contrast to passing arguments by *value* as discussed in Chap. 7.

When an argument is passed by value, the data item is *copied* to the function. Thus, any alteration made to the data item within the function is not carried over into the calling routine (see Sec. 7.5). When an argument is passed by reference, however (i.e., when a pointer is passed to a function), the *address* of a data item is passed to the function. The contents of that address can be accessed freely, either within the function or within the calling routine. Moreover, any change that is made to the data item (i.e., to the contents of the address) will be recognized in both the function and the calling routine. Thus, the use of a pointer as a function argument permits the corresponding data item to be altered globally from within the function.

When pointers are used as arguments to a function, some care is required with the formal argument declarations within the function. Specifically, formal pointer arguments that must each be preceded by an asterisk. Function prototypes are written in the same manner. If a function declaration does not include variable names, the data type of each pointer argument must be followed by an asterisk. The use of pointer arguments is illustrated in the following example.

EXAMPLE 10.7 Here is a simple C program that illustrates the difference between ordinary arguments, which are passed by value, and pointer arguments, which are passed by reference.

```

#include <stdio.h>

void funct1(int u, int v);           /* function prototype */
void funct2(int *pu, int *pv);       /* function prototype */

main()
{
    int u = 1;
    int v = 3;

    printf("\nBefore calling funct1: u=%d    v=%d", u, v);
    funct1(u, v);
    printf("\nAfter calling funct1:   u=%d    v=%d", u, v);

    printf("\n\nBefore calling funct2:   u=%d    v=%d", u, v);
    funct2(&u, &v);
    printf("\nAfter calling funct2:   u=%d    v=%d", u, v);
}

void funct1(int u, int v)
{
    u = 0;
    v = 0;
    printf("\nWithin funct1:           u=%d    v=%d", u, v);
    return;
}

void funct2(int *pu, int *pv)
{
    *pu = 0;
    *pv = 0;
    printf("\nWithin funct2:           *pu=%d *pv=%d", *pu, *pv);
    return;
}

```

This program contains two functions, called `funct1` and `funct2`. The first function, `funct1`, receives two integer variables as arguments. These variables are originally assigned the values 1 and 3, respectively. The values are then changed, to 0, 0 within `funct1`. The new values are not recognized in `main`, however, because the arguments were passed by value, and any changes to the arguments are local to the function in which the changes occur.

Now consider the second function, `funct2`. This function receives two *pointers* to integer variables as its arguments. The arguments are identified as pointers by the indirection operators (i.e., the asterisks) that appear in the argument declaration. In addition, the argument declaration indicates that the pointers contain the addresses of *integer* quantities.

Within `funct2`, the contents of the pointer addresses are reassigned the values 0, 0. Since the addresses are recognized in both `funct2` and `main`, the reassigned values will be recognized within `main` after the call to `funct2`. Therefore, the integer variables `u` and `v` will have their values changed from 1, 3 to 0, 0.

The six `printf` statements illustrate the values of `u` and `v`, and their associated values `*pu` and `*pv`, within `main` and within the two functions. Hence, the following output is generated when the program is executed.

```

Before calling funct1: u=1    v=3
Within funct1:         u=0    v=0
After calling funct1:  u=1    v=3

```

```

Before calling funct2: u=1  v=3
Within funct2:      *pu=0 *pv=0
After calling funct2: u=0  v=0

```

Notice that the values of *u* and *v* are unchanged within *main* after the call to *funct1*, though the values of these variables are changed within *main* after the call to *funct2*. Thus, the output illustrates the local nature of the alterations within *funct1*, and the global nature of the alterations within *funct2*.

This example contains some additional features that should be pointed out. Notice, for example, the function prototype

```
void funct2(int *pu, int *pv);
```

The items in parentheses identify the arguments as pointers to integer quantities. The pointer variables, *pu* and *pv*, have not been declared elsewhere in *main*. This is permitted in the function prototype, however, because *pu* and *pv* are dummy arguments rather than actual arguments. The function declaration could also have been written without any variable names, as

```
void funct2(int *, int *);
```

Now consider the declaration of the formal arguments within the first line of *funct2*, i.e.,

```
void funct2(int *pu, int *pv)
```

The formal arguments *pu* and *pv* are consistent with the dummy arguments in the function prototype. In this example the corresponding variable names are the same, though this is generally not required.

Finally, notice the manner in which *u* and *v* are accessed within *funct2*, i.e.,

```
*pu = 0;
*pv = 0;
```

Thus, *u* and *v* are accessed indirectly, by referencing the contents of the addresses represented by the pointers *pu* and *pv*. This is necessary because the variables *u* and *v* are not recognized as such within *funct2*.

We have already mentioned the fact that an array name is actually a pointer to the array; i.e., the array name represents the address of the first element in the array (see Sec. 9.3). Therefore, an array name is treated as a pointer when it is passed to a function. However, it is not necessary to precede the array name with an ampersand within the function call.

An array name that appears as a formal argument within a function definition can be declared either as a pointer or as an array of unspecified size, as shown in Sec. 9.3. The choice is a matter of personal preference, though it will often be determined by the manner in which the individual array elements are accessed within the function (more about this in the next section).

EXAMPLE 10.8 Analyzing a Line of Text Suppose we wish to analyze a line of text by examining each of the characters and determining into which of several different categories it falls. In particular, suppose we count the number of vowels, consonants, digits, whitespace characters and “other” characters (punctuation, operators, brackets, etc.) This can easily be accomplished by reading in a line of text, storing it in a one-dimensional character array, and then analyzing the individual array elements. An appropriate counter will be incremented for each character. The value of each counter (number of vowels, number of consonants, etc.) can then be written out after all of the characters have been analyzed.

Let us write a complete C program that will carry out such an analysis. To do so, we first define the following symbols.

```

line = an 80-element character array containing the line of text
vowels = an integer counter indicating the number of vowels
consonants = an integer counter indicating the number of consonants
digits = an integer counter indicating the number of digits

```

whitespc = an integer counter indicating the number of whitespace characters (blank spaces or tabs)

other = an integer counter indicating the number of characters that do not fall into any of the preceding categories

Notice that newline characters are not included in the "whitespace" category, because there can be no newline characters within a single line of text.

We will structure the program so that the line of text is read into the main portion of the program, and then passed to a function where it will be analyzed. The function will return the value of each counter after all of the characters have been analyzed. The results of the analysis (i.e., the value of each counter) will then be displayed from the main portion of the program.

The actual analysis can be carried out by creating a loop to examine each of the characters. Within the loop we first convert each character that is a letter to uppercase. This avoids the need to distinguish between uppercase and lowercase letters. We can then categorize the character using a nest of if - else statements. Once the proper category has been identified, the corresponding counter is incremented. The entire process is repeated until the string termination character (\0) has been found.

The complete C program is shown below.

```

/* count the number of vowels, consonants, digits, whitespace characters,
   and 'other' characters in a line of text */
```

```

#include <stdio.h>
#include <ctype.h>

/* function prototype */
void scan_line(char line[], int *pv, int *pc, int *pd, int *pw, int *po);

main()
{
    char line[80];           /* line of text */
    int vowels = 0;          /* vowel counter */
    int consonants = 0;      /* consonant counter */
    int digits = 0;          /* digit counter */
    int whitespc = 0;        /* whitespace counter */
    int other = 0;           /* remaining character counter */

    printf("Enter a line of text below:\n");
    scanf("%[^\\n]", line);

    scan_line(line, &vowels, &consonants, &digits, &whitespc, &other);

    printf("\nNo. of vowels: %d", vowels);
    printf("\nNo. of consonants: %d", consonants);
    printf("\nNo. of digits: %d", digits);
    printf("\nNo. of whitespace characters: %d", whitespc);
    printf("\nNo. of other characters: %d", other);
}

void scan_line(char line[], int *pv, int *pc, int *pd, int *pw, int *po)
/* analyze the characters in a line of text */
{
    char c;                 /* uppercase character */
    int count = 0;           /* character counter */

```

```

while ((c = toupper(line[count])) != '\0') {
    if (c == 'A' || c == 'E' || c == 'I' || c == 'O' || c == 'U')
        ++ *pv;                                /* vowel */
    else if (c >= 'A' && c <= 'Z')
        ++ *pc;                                /* consonant */
    else if (c >= '0' && c <= '9')
        ++ *pd;                                /* digit */
    else if (c == ' ' || c == '\t')
        ++ *pw;                                /* whitespace */
    else
        ++ *po;                                /* other */
    ++count;
}
return;
}

```

Notice the function prototype for `scan_line` that appears at the beginning of the program. In particular, notice the use of the `void` data type, and notice the manner in which the argument data types are specified. Note the distinction between the array argument and the remaining pointer arguments.

Also, observe the manner in which the actual arguments are written in the call to `scan_line`. The array argument, `line`, is not preceded by an ampersand, since arrays are, by definition, pointers. Each of the remaining arguments must be preceded by an ampersand so that its address, rather than its value, is passed to the function.

Now consider the function `scan_line`. All of the formal arguments, including `line`, are pointers. However, `line` is declared as an array whose size is unspecified, whereas the remaining arguments are specifically declared as pointers. It is possible (and quite common) to declare `line` as a pointer rather than an array. Thus, the first line of `scan_line` could have been written as

```
void scan_line(char *line, int *pv, int *pc, int *pd, int *pw, int *po)
```

rather than as shown in the program listing. To be consistent, the corresponding function prototype would then be written in a similar manner.

Incrementing the various counters also requires some explanation. First, note that it is the *content* of each address (i.e., the *object* of each pointer) that is incremented. Second, note that each indirection expression (e.g., `*pv`) is *preceded* by the unary operator `++`. Since the unary operators are evaluated from right to left, we are assured that the content of each address, rather than the address itself, is increased in value.

Here is a typical dialog that might be encountered when the program is executed. (The line of text entered by the user is underlined.)

Enter a line of text below:

Personal computers with memories in excess of 4096 KB are now quite common.

The corresponding output is:

```
No. of vowels: 23
No. of consonants: 35
No. of digits: 4
No. of whitespace characters: 12
No. of other characters: 1
```

Thus, we see that this particular line of text contains 23 vowels, 35 consonants, 4 digits, 12 whitespace characters (blank spaces), and one other character (the period).

Recall that the `scanf` function requires those arguments that are ordinary variables to be preceded by ampersands (see Sec. 4.4). However, array names are exempted from this requirement. This may have seemed somewhat mysterious back in Chap. 4 but it should now make sense, considering what we now know about array names and addresses. Thus, the `scanf` function requires that the *addresses* of the data items being entered into the computer's memory be specified. The ampersands provide a means for accessing the addresses of ordinary single-valued variables. Ampersands are not required with array names, since array names themselves represent addresses.

EXAMPLE 10.9 The skeletal structure of a C program is shown below (repeated from Example 4.5).

```
#include <stdio.h>

main()
{
    char item[20];
    int partno;
    float cost;

    . . .

    scanf("%s %d %f", item, &partno, &cost);

    . . .
}
```

The `scanf` statement causes a character string, an integer quantity and a floating-point quantity to be entered into the computer and stored in the addresses associated with `item`, `partno` and `cost`, respectively. Since `item` is the name of an array, it is understood to represent an address. Hence, `item` need not (cannot) be preceded by an ampersand within the `scanf` statement. On the other hand, `partno` and `cost` are conventional variables. Therefore they must be written as `&partno` and `&cost` within the `scanf` statement. The ampersands are required in order to access the *addresses* of these variables rather than their values.

If the `scanf` function is used to enter a single array element rather than an entire array, the name of the array element must be preceded by an ampersand, as shown below (from Example 9.8):

```
scanf("%f", &list[count]);
```

It is possible to pass a *portion* of an array, rather than an entire array, to a function. To do so, the address of the first array element to be passed must be specified as an argument. The remainder of the array, starting with the specified array element, will then be passed to the function.

EXAMPLE 10.10 The skeletal structure of a C program is shown below.

```
#include <stdio.h>

void process(float z[]);

main()
{
    float z[100];
    . . .
    /* enter values for elements of z */
    . . .
    process(&z[50]);
    . . .
}
```

```

void process(float f[])
{
    . . .
    /* process elements of f */
    . . .
    return;
}

```

Within **main**, **z** is declared to be a 100-element, floating-point array. After the elements of **z** are entered into the computer, the address of **z[50]** (i.e., **&z[50]**) is passed to the function **process**. Hence, the last 50 elements of **z** (i.e., the elements **z[50]** through **z[99]**) will be available to **process**.

In the next section we will see that the address of **z[50]** can be written as **z + 50** rather than **&z[50]**. Therefore, the call to **process** can appear as **process(z + 50)** rather than **process(&z[50])**, as shown above. Either method may be used, depending on the programmer's preferences.

Within **process**, the corresponding array is referred to as **f**. This array is declared to be a floating-point array whose size is unspecified. Thus, the fact that the function receives only a portion of **z** is immaterial; if all of the array elements are altered within **process**, only the last 50 elements will be affected within **main**.

Within **process**, it may be desirable to declare the formal argument **f** as a pointer to a floating-point quantity rather than an array name. Thus, the outline of **process** may be written as

```

void process(float *f)
{
    . . .
    /* process elements of f */
    . . .
    return;
}

```

Notice the difference between the formal argument declarations in the two function outlines. Both declarations are valid.

A function can also return a pointer to the calling portion of the program. To do so, the function definition and any corresponding function declarations must indicate that the function will return a pointer. This is accomplished by preceding the function name by an asterisk. The asterisk must appear in both the function definition and the function declarations.

EXAMPLE 10.11 Shown below is the skeletal structure of a C program that transfers a double-precision array to a function and returns a pointer to one of the array elements.

```

#include <stdio.h>

double *scan(double z[]);

main()
{
    double z[100];           /* array declaration */
    double *pz;              /* pointer declaration */

    /* enter values for elements of z */

    . . .

    pz = scan(z);

    . . .
}

```

```

    double *scan(double f[])
{
    double *pf;           /* pointer declaration */

    . . .

/* process elements of f */

pf = . . . .;

return(pf);
}

```

Within `main` we see that `z` is declared to be a 100-element, double-precision array, and `pz` is a pointer to a double-precision quantity. We also see a declaration for the function `scan`. Note that `scan` will accept a double-precision array as an argument, and it will return a pointer to (i.e., the address of) a double-precision quantity. The asterisk preceding the function name (`*scan`) indicates that the function will return a pointer.

Within the function definition, the first line indicates that `scan` accepts one formal parameter (`f[]`) and returns a pointer to a double-precision quantity. The formal parameter will be a one-dimensional, double-precision array. The outline suggests that the address of one of the array elements is assigned to the pointer `pf` during or after the processing of the array elements. This address is then returned to `main`, where it is assigned to the pointer variable `pz`.

10.4 POINTERS AND ONE-DIMENSIONAL ARRAYS

Recall that an array name is really a pointer to the first element in the array. Therefore, if `x` is a one-dimensional array, then the address of the first array element can be expressed as either `&x[0]` or simply as `x`. Moreover, the address of the second array element can be written as either `&x[1]` or as `(x + 1)`, and so on. In general, the address of array element $(i + 1)$ can be expressed as either `&x[i]` or as `(x + i)`. Thus we have two different ways to write the address of any array element: We can write the actual array element, preceded by an ampersand; or we can write an expression in which the subscript is added to the array name.

In the latter case, it should be understood that we are dealing with a very special and unusual type of expression. In the expression `(x + i)`, for example, `x` represents an address, whereas `i` represents an integer quantity. Moreover, `x` is the name of an array whose elements may be characters, integers, floating-point quantities, etc. (though all of the the array elements must be of the same data type). Thus, we are not simply adding numerical values. Rather, we are specifying an address that is a certain number of memory cells beyond the address of the first array element. Or, in simpler terms, we are specifying a location that is `i` array elements beyond the first. Hence, the expression `(x + i)` is a symbolic representation for an address specification rather than an arithmetic expression.

Recall that the number of memory cells associated with an array element will depend upon the data type of the array as well as the particular computer's architecture. With some computers, for example, an integer quantity occupies two bytes (two memory cells), a floating-point quantity requires four bytes, and a double-precision quantity requires eight bytes of memory. With other computers, an integer quantity may require four bytes, and floating-point and double-precision quantities may each require eight bytes. And so on.

When writing the address of an array element in the form `(x + i)`, however, you need not be concerned with the number of memory cells associated with each type of array element; the C compiler adjusts for this automatically. You must specify only the address of the first array element (i.e., the name of the array) and the number of array elements beyond the first (i.e., a value for the subscript). The value of `i` is sometimes referred to as an *offset* when used in this manner.

Since `&x[i]` and `(x + i)` both represent the address of the `i`th element of `x`, it would seem reasonable that `x[i]` and `*(x + i)` both represent the contents of that address, i.e., the *value* of the `i`th element of `x`. This is indeed the case. The two terms are interchangeable. Hence, either term can be used in any particular application. The choice depends upon your individual preferences.

EXAMPLE 10.12 Here is a simple program that illustrates the relationship between array elements and their addresses.

```
#include <stdio.h>

main()
{
    static int x[10] = {10, 11, 12, 13, 14, 15, 16, 17, 18, 19};
    int i;

    for (i = 0; i <= 9; ++i) {
        /* display an array element */
        printf("\ni= %d      x[i]= %d      *(x+i)= %d", i, x[i], *(x+i));

        /* display the corresponding array address */
        printf("      &x[i]= %X      x+i= %X", &x[i], (x+i));
    }
}
```

This program defines a one-dimensional, 10-element integer array *x*, whose elements are assigned the values 10, 11, . . . , 19. The action portion of the program consists of a loop that displays the value and the corresponding address of each array element. Note that the value of each array element is specified in two different ways, as *x[i]* and as **(x+i)*, in order to illustrate their equivalence. Similarly, the address of each array element is specified in two different ways, as *&x[i]* and as *(x+i)*, for the same reason. Therefore the value and the address of each array element should appear twice.

Execution of this program results in the following output.

i= 0	x[i]= 10	*(x+i)= 10	&x[i]= 72	x+i= 72
i= 1	x[i]= 11	*(x+i)= 11	&x[i]= 74	x+i= 74
i= 2	x[i]= 12	*(x+i)= 12	&x[i]= 76	x+i= 76
i= 3	x[i]= 13	*(x+i)= 13	&x[i]= 78	x+i= 78
i= 4	x[i]= 14	*(x+i)= 14	&x[i]= 7A	x+i= 7A
i= 5	x[i]= 15	*(x+i)= 15	&x[i]= 7C	x+i= 7C
i= 6	x[i]= 16	*(x+i)= 16	&x[i]= 7E	x+i= 7E
i= 7	x[i]= 17	*(x+i)= 17	&x[i]= 80	x+i= 80
i= 8	x[i]= 18	*(x+i)= 18	&x[i]= 82	x+i= 82
i= 9	x[i]= 19	*(x+i)= 19	&x[i]= 84	x+i= 84

The output clearly illustrates the distinction between *x[i]*, which represents the value of the *i*th array element, and *&x[i]*, which represents its address. Moreover, we see that the value of the *i*th array element can be represented by either *x[i]* or **(x+i)*, and the address of the *i*th element can be represented by either *&x[i]* or *x+i*. Thus we see another comparison, between **(x+i)*, which also represents the value of the *i*th element, and *x+i*, which also represents its address.

Notice, for example, that the first array element (corresponding to *i* = 0) has been assigned a value of 10 and a (hexadecimal) address of 72. The second array element has a value of 11 and an address of 74, etc. Thus, memory location 72 will contain the integer value 10, location 74 will contain 11, and so on.

You should understand that these addresses are assigned automatically by the compiler.

When assigning a *value* to an array element such as *x[i]*, the left side of the assignment statement may be written as either *x[i]* or as **(x + i)*. Thus, a value may be assigned directly to an array element, or it may be assigned to the memory area whose address is that of the array element.

On the other hand, it is sometimes necessary to assign an *address* to an identifier. In such situations, a pointer variable must appear on the left side of the assignment statement. It is *not* possible to assign an arbitrary address to an array name or to an array element. Thus, expressions such as *x*, *(x + i)* and *&x[i]* cannot appear on the left side of an assignment statement. Moreover, the address of an array cannot arbitrarily be altered, so that expressions such as *++x* are not permitted.

EXAMPLE 10.13 Consider the skeletal structure of the C program shown below.

```
#include <stdio.h>

main()
{
    int line[80];
    int *p1;

    . . . .

    /* assign values */
    line[2] = line[1];
    line[2] = *(line + 1);
    *(line + 2) = line[1];
    *(line + 2) = *(line + 1);

    /* assign addresses */
    p1 = &line[1];
    p1 = line + 1;
}
```

Each of the first four assignment statements assigns the value of the second array element (i.e., `line[1]`) to the third array element (`line[2]`). Thus, the four statements are all equivalent. An experienced programmer would probably choose either the first or the fourth, however, in order that the notation be consistent.

The last two assignment statements each assigns the *address* of the second array element to the pointer `p1`. We might choose to do this in an actual program if it were necessary to “tag” the address of `line[1]` for some reason.

Note that *the address of one array element cannot be assigned to some other array element*. Thus we *cannot* write a statement such as

```
&line[2] = &line[1];
```

On the other hand, we *can* assign the *value* of one array element to another through a pointer if we wish, e.g.,

```
p1 = &line[1];
line[2] = *p1;
```

or

```
p1 = line + 1;
*(line + 2) = *p1;
```

If a numerical array is defined as a pointer variable, the array elements cannot be assigned initial values. Therefore, a conventional array definition is required if initial values will be assigned to the elements of a numerical array. However, a *character-type* pointer variable can be assigned an entire string as a part of the variable declaration. Thus, a string can conveniently be represented by either a one-dimensional character array or a character pointer.

EXAMPLE 10.14 Shown below is a simple C program in which two strings are represented as one-dimensional character arrays.

```
#include <stdio.h>

char x[] = "This string is declared externally\n\n";
```

```
main()
{
    static char y[] = "This string is declared within main";
    printf("%s", x);
    printf("%s", y);
}
```

The first string is assigned to the external array `x[]`. The second string is assigned to the static array `y[]`, which is defined within `main`. This second definition occurs within a function; therefore `y[]` must be defined as a `static` array so that it can be initialized.

Here is a different version of the same program. The strings are now assigned to pointer variables rather than to one-dimensional arrays.

```
#include <stdio.h>

char *x = "This string is declared externally\n\n";

main()
{
    char *y = "This string is declared within main";
    printf("%s", x);
    printf("%s", y);
}
```

The external pointer variable `x` points to the beginning of the first string, whereas the pointer variable `y`, declared within `main`, points to the beginning of the second string. Note that `y` can now be initialized without being declared `static`.

Execution of either program produces the following output.

```
This string is declared externally
This string is declared within main
```

Syntactically, of course, it is possible to declare a pointer variable `static`. However, there is no reason to do so in this example.

10.5 DYNAMIC MEMORY ALLOCATION

Since an array name is actually a pointer to the first element within the array, it should be possible to define the array as a pointer variable rather than as a conventional array. Syntactically, the two definitions are equivalent. However, a conventional array definition results in a fixed block of memory being reserved at the beginning of program execution, whereas this does not occur if the array is represented in terms of a pointer variable. Therefore, the use of a pointer variable to represent an array requires some type of initial memory assignment before the array elements are processed. This is known as *dynamic memory allocation*. Generally, the `malloc` library function is used for this purpose, as illustrated in the next example.

EXAMPLE 10.15 Suppose `x` is a one-dimensional, 10-element array of integers. It is possible to define `x` as a pointer variable rather than an array. Thus, we can write

```
int *x;
```

rather than

```
int x[10];
```

or

```
#define SIZE 10
int x[SIZE];
```

However, *x* is not automatically assigned a memory block when it is defined as a pointer variable, though a block of memory large enough to store 10 integer quantities will be reserved in advance when *x* is defined as an array.

To assign sufficient memory for *x*, we can make use of the library function `malloc`, as follows.

```
x = (int *) malloc(10 * sizeof(int));
```

This function reserves a block of memory whose size (in bytes) is equivalent to 10 integer quantities. As written, the function returns a pointer to an integer. This pointer indicates the beginning of the memory block. In general, the type cast preceding `malloc` must be consistent with the data type of the pointer variable. Thus, if *y* were defined as a pointer to a double-precision quantity and we wanted enough memory to store 10 double-precision quantities, we would write

```
y = (double *) malloc(10 * sizeof(double));
```

If the declaration is to include the assignment of initial values, however, then x must be defined as an array rather than a pointer variable. For example,

```
int x[10] = {1, 2, 3, 4, 5, 6, 7, 8, 9, 10};
```

or

```
int x[] = {1, 2, 3, 4, 5, 6, 7, 8, 9, 10};
```

When programming in C, it is not unusual to use pointer expressions rather than references to individual array elements. The resulting programs may appear strange at first, though they are straightforward once you become comfortable accessing values that are stored in specified addresses. Generally, a small amount of practice is all that is required.

EXAMPLE 10.16 Reordering a List of Numbers To illustrate the use of pointers with dynamic memory allocation, let us once again consider the problem of reordering a list of integers, as described in Example 9.13. Now, however, we will utilize pointer expressions to access numerical values rather than refer explicitly to individual array elements. In all other respects, we present a program that is identical to that given in Example 9.13.

Here is the complete C program.

```
/* reorder a one-dimensional, integer array from smallest to largest,
   using pointer notation */
```

```
#include <stdio.h>
#include <stdlib.h>

void reorder(int n, int *x);

main()
{
    int i, n, *x;

    /* read in a value for n */
    printf("\nHow many numbers will be entered? ");
    scanf("%d", &n);
    printf("\n");

    /* allocate memory */
    x = (int *) malloc(n * sizeof(int));
```

```

/* read in the list of numbers */
for (i = 0; i < n; ++i)  {
    printf("i = %d  x = ", i + 1);
    scanf("%d", x + i);
}

/* reorder all array elements */
reorder(n, x);

/* display the reordered list of numbers */
printf("\n\nReordered List of Numbers:\n\n");
for (i = 0; i < n; ++i)
    printf("i = %d  x = %d\n", i + 1, *(x + i));
}

void reorder(int n, int *x) /* rearrange the list of numbers */
{
    int i, item, temp;

    for (item = 0; item < n - 1; ++item)

        /* find the smallest of all remaining elements */
        for (i = item + 1; i < n; ++i)

            if (*(x + i) < *(x + item))  {

                /* interchange two elements */
                temp = *(x + item);
                *(x + item) = *(x + i);
                *(x + i) = temp;
            }
    return;
}

```

In this program, the integer array is defined as a pointer to an integer. Memory is initially assigned to the pointer variable via the `malloc` library function. Elsewhere in the program, pointer notation is used to process the individual array elements. For example, the function prototype now specifies that the second argument is a pointer to an integer quantity rather than an integer array. This pointer will identify the beginning of the integer array.

We also see that the `scanf` function now specifies the address of the *i*th element as `x + i` rather than `&x[i]`. Similarly, the `printf` function now represents the value of the *i*th element as `*(x + i)` rather than `x[i]`. The call to `reorder`, however, is the same as in the earlier program—namely, `reorder(n, x);`.

Within the function `reorder`, we see that the second formal argument is now defined as a pointer variable rather than an integer array. This is consistent with the function prototype. Even more pronounced, however, are the differences in the `if` statement. In particular, notice that each reference to an array element is now written as the contents of an address. Thus `x[i]` is now written as `*(x + i)`, and `x[item]` is now written as `*(x + item)`. This compound `if` statement can be viewed as a conditional interchange involving the contents of two different addresses, rather than an interchange of two different elements within a conventional array.

You should compare this program with that shown in Example 9.13 in order to appreciate the differences. Both programs will generate identical results with the same input data. However, you should understand the syntactic differences between the two programs.

An important advantage of dynamic memory allocation is the ability to reserve as much memory as may be required during program execution, and then release this memory when it is no longer needed. Moreover, this process may be repeated many times during execution of a program. The library functions `malloc` and `free` are used for these purposes, as illustrated in Example 11.32 (see Sec. 11.6).

10.6 OPERATIONS ON POINTERS

In Sec. 10.4 we saw that an integer value can be added to an array name in order to access an individual array element. The integer value is interpreted as an array subscript; it represents the location of the desired array element relative to the first element in the array. This works because all of the array elements are of the same data type, and therefore each array element occupies the same number of memory cells (i.e., the same number of bytes or words). The actual number of memory cells that separate the two array elements will depend on the data type of the array, though this is taken care of automatically by the compiler and therefore need not concern you directly.

This concept can be extended to pointer variables in general. Thus, an integer value can be added to or subtracted from a pointer variable, though the resulting expression must be interpreted very carefully. Suppose, for example, that px is a pointer variable that represents the address of some variable x . We can write expressions such as $++\text{px}$, $-\text{px}$, $(\text{px} + 3)$, $(\text{px} + i)$, and $(\text{px} - i)$, where i is an integer variable. Each expression will represent an address that is located some distance from the original address represented by px . The exact distance will be the product of the integer quantity and the number of bytes or words associated with the data item to which px points. Suppose, for example, that px points to an integer quantity, and each integer quantity requires two bytes of memory. Then the expression $(\text{px} + 3)$ will result in an address that is six bytes beyond the integer to which px points, as illustrated in Fig. 10.3. It should be understood, however, that this new address will *not* necessarily represent the address of another data item, particularly if the data items stored between the two addresses involve different data types.

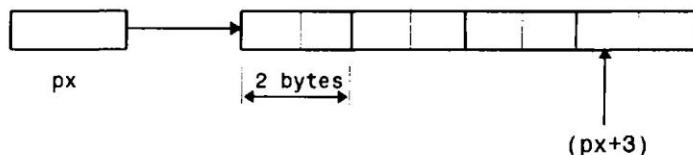


Fig. 10.3

EXAMPLE 10.17 Consider the simple C program shown below.

```
#include <stdio.h>

main()
{
    int *px;           /* pointer to an integer */
    int i = 1;
    float f = 0.3;
    double d = 0.005;
    char c = '*';

    px = &i;
    printf("Values: i=%i f=%f d=%f c=%c\n\n", i, f, d, c);
    printf("Addresses: &i=%X &f=%X &d=%X &c=%X\n\n", &i, &f, &d, &c);
    printf("Pointer values: px=%X px + 1=%X px + 2=%X px + 3=%X",
           px, px + 1, px + 2, px + 3);
}
```

This program displays the values and addresses associated with four different types of variables: i , an integer variable; f , a floating-point variable; d , a double-precision variable; and c , a character variable. The program also makes use of a pointer variable, px , which represents the address of i . The values of px , $\text{px} + 1$, $\text{px} + 2$ and $\text{px} + 3$ are also displayed, so that they may be compared with the addresses of the different variables.

Execution of the program results in the following output.

```
Values:   i=1    f=0.300000  d=0.005000  c=*
Addresses: &i=117E  &f=1180  &d=1186  &c=118E
Pointer values: px=117E  px + 1=1180  px + 2=1182  px + 3=1184
```

The first line simply displays the values of the variables, and the second line displays their addresses, as assigned by the compiler. Notice that the number of bytes associated with each data item is different. Thus, the integer value represented by *i* requires two bytes (specifically, addresses 117E and 117F). The floating-point value represented by *f* appears to be assigned six bytes (addresses 1180 through 1185), though only four bytes (addresses 1180 through 1183) are actually used for this purpose. (Compilers allocate memory space according to their own rules.) However, eight bytes are required for the double-precision value represented by *d* (addresses 1186 through 118D). And finally, the character represented by *c* begins in address 118E. Only one byte is required to store this single character, though the output does not indicate the number of bytes between this character and the next data item.

Now consider the third line of output, which contains the addresses represented by the pointer expressions. Clearly, *px* represents the address of *i* (i.e., 117E). This comes as no surprise, since this address was explicitly assigned to *px* by the expression *px = &i*. However, *px + 1* moves over only two bytes, to 1180, and *px + 2* moves over another two bytes, to 1182, and so on. The reason is that *px* points to an integer quantity, and integer quantities each require two bytes with this particular C compiler. Therefore, when integer constants are added to *px*, the constants are interpreted in terms of two-byte multiples.

If *px* is defined as a pointer to a different type of object (e.g., a character or a floating-point quantity), then any integer constant that is added to or subtracted from the pointer will be interpreted differently. In particular, each integer value will represent an equivalent number of individual bytes if *px* points to a character, or a corresponding number of four-byte multiples if *px* points to a floating-point quantity. You are encouraged to verify this on your own.

One pointer variable can be subtracted from another provided both variables point to elements of the same array. The resulting value indicates the number of words or bytes separating the corresponding array elements.

EXAMPLE 10.18 In the program shown below, two different pointer variables point to the first and the last elements of an integer array.

```
#include <stdio.h>

main()
{
    int *px, *py;          /* integer pointers */
    static int a[6] = {1, 2, 3, 4, 5, 6};

    px = &a[0];
    py = &a[5];
    printf("px=%X  py=%X", px, py);
    printf("\n\npy - px=%X", py - px);
}
```

In particular, the pointer variable *px* points to *a[0]*, and *py* points to *a[5]*. The difference, *py - px*, should be 5, since *a[5]* is the fifth element beyond *a[0]*.

Execution of the program results in the following output.

```
px=52  py=5C
py - px=5
```

The first line indicates that the address of *a[0]* is 52, and the address of *a[5]* is 5C. The difference between these two hexadecimal numbers is 10 (when converted to decimal). Thus, *a[5]* is stored at an address which is 10 bytes beyond the

address of `a[0]`. Since each integer quantity occupies two bytes, we would expect the difference between `py` and `px` to be $10/2 = 5$. The second line of output confirms this value.

Pointer variables can be compared provided both variables are of the same data type. Such comparisons can be useful when both pointer variables point to elements of the same array. The comparisons can test for either equality or inequality. Moreover, a pointer variable can be compared with zero (which is usually expressed as `NULL` when used in this manner, as explained in Sec. 10.2).

EXAMPLE 10.19 Suppose `px` and `py` are pointer variables that point to elements of the same array. Several logical expressions involving these two variables are shown below. All of the expressions are syntactically correct.

```
(px < py)  
(px >= py)  
(px == py)  
(px != py)  
(px == NULL)
```

These expressions can be used as any other logical expression. For example,

```
if (px < py)  
    printf("px < py");  
else  
    printf("px >= py");
```

Expressions such as `(px < py)` indicate whether or not the element associated with `px` is ranked ahead of the element associated with `py` (i.e., whether or not the subscript associated with `*px` is less than the subscript associated with `*py`).

You should understand that the operations discussed previously are the *only* operations that can be carried out on pointers. These permissible operations are summarized below.

1. A pointer variable can be assigned the address of an ordinary variable (e.g., `pv = &v`).
2. A pointer variable can be assigned the value of another pointer variable (e.g., `pv = px`) provided both pointers point to objects of the same data type .
3. A pointer variable can be assigned a null (zero) value (e.g., `pv = NULL`, where `NULL` is a symbolic constant that represents the value 0).
4. An integer quantity can be added to or subtracted from a pointer variable (e.g., `pv + 3`, `++pv`, etc.)
5. One pointer variable can be subtracted from another provided both pointers point to elements of the same array.
6. Two pointer variables can be compared provided both pointers point to objects of the same data type.

Other arithmetic operations on pointers are not allowed. Thus, a pointer variable cannot be multiplied by a constant; two pointer variables cannot be added; and so on. Also, you are again reminded that an ordinary variable cannot be assigned an arbitrary address (i.e., an expression such as `&x` cannot appear on the left side of an assignment statement).

10.7 POINTERS AND MULTIDIMENSIONAL ARRAYS

Since a one-dimensional array can be represented in terms of a pointer (the array name) and an offset (the subscript), it is reasonable to expect that a multidimensional array can also be represented with an equivalent pointer notation. This is indeed the case. A two-dimensional array, for example, is actually a collection of

one-dimensional arrays. Therefore, we can define a two-dimensional array as a pointer to a group of contiguous one-dimensional arrays. Thus, a two-dimensional array declaration can be written as

```
data-type (*ptvar)[expression 2];
```

rather than

```
data-type array[expression 1][expression 2];
```

This concept can be generalized to higher-dimensional arrays; that is,

```
data-type (*ptvar)[expression 2][expression 3] . . . [expression n];
```

replaces

```
data-type array[expression 1][expression 2] . . . [expression n];
```

In these declarations *data-type* refers to the data type of the array, *ptvar* is the name of the pointer variable, *array* is the corresponding array name, and *expression 1*, *expression 2*, . . . , *expression n* are positive-valued integer expressions that indicate the maximum number of array elements associated with each subscript.

Notice the parentheses that surround the array name and the preceding asterisk in the pointer version of each declaration. These parentheses must be present. Without them we would be defining an array of pointers rather than a pointer to a group of arrays, since these particular symbols (i.e., the square brackets and the asterisk) would normally be evaluated right to left. We will say more about this in the next section.

EXAMPLE 10.20 Suppose *x* is a two-dimensional integer array having 10 rows and 20 columns. We can declare *x* as

```
int (*x)[20];
```

rather than

```
int x[10][20];
```

In the first declaration, *x* is defined to be a pointer to a group of contiguous, one-dimensional, 20-element integer arrays. Thus, *x* points to the first 20-element array, which is actually the first row (i.e., row 0) of the original two-dimensional array. Similarly, (*x* + 1) points to the second 20-element array, which is the second row (row 1) of the original two-dimensional array, and so on, as illustrated in Fig. 10.4.

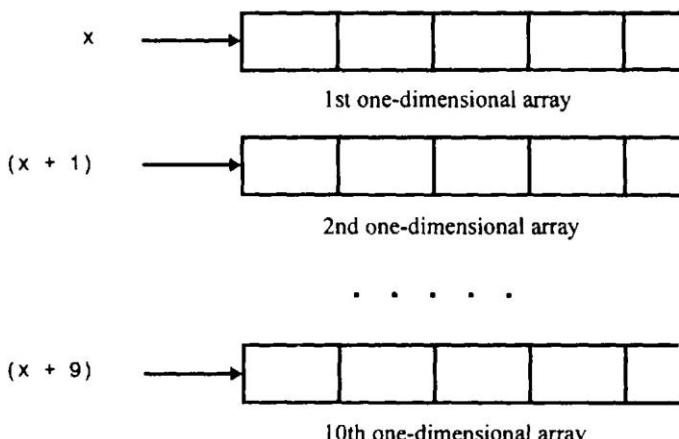


Fig. 10.4

Now consider a three-dimensional floating-point array t . This array can be defined as

```
float (*t)[20][30];
```

rather than

```
float t[10][20][30];
```

In the first declaration, t is defined as a pointer to a group of contiguous two-dimensional, 20×30 floating-point arrays. Hence, t points to the first 20×30 array, $(t + 1)$ points to the second 20×30 array, and so on.

An individual array element within a multidimensional array can be accessed by the repeated use of the indirection operator. Usually, however, this procedure is more awkward than the conventional method for accessing an array element. The following example illustrates the use of the indirection operator.

EXAMPLE 10.21 Suppose x is a two-dimensional integer array having 10 rows and 20 columns, as declared in the previous example. The item in row 2, column 5 can be accessed by writing either

```
x[2][5]
```

or

```
*(*(x + 2) + 5)
```

The second form requires some explanation. First, note that $(x + 2)$ is a pointer to row 2. Therefore the object of this pointer, $*(x + 2)$, refers to the entire row. Since row 2 is a one-dimensional array, $*(x + 2)$ is actually a pointer to the first element in row 2. We now add 5 to this pointer. Hence, $(*x + 2) + 5$ is a pointer to element 5 (i.e., the sixth element) in row 2. The object of this pointer, $*(x + 2) + 5$, therefore refers to the item in column 5 of row 2, which is $x[2][5]$. These relationships are illustrated in Fig. 10.5.

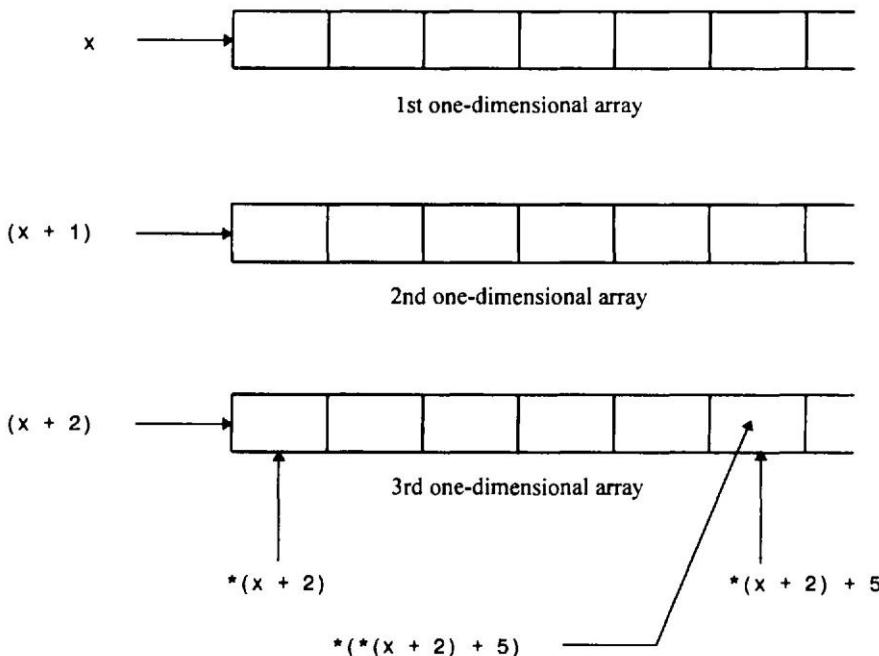


Fig. 10.5

Programs that make use of multidimensional arrays can be written in several different ways. In particular, there are different ways to define the arrays, and different ways to process the individual array elements. The choice of one method over another is often a matter of personal preference. In applications involving numerical arrays, it is often easier to define the arrays in the conventional manner, thus avoiding any possible subtleties associated with initial memory assignments. The following example, however, illustrates the use of pointer notation to process multidimensional numerical arrays.

EXAMPLE 10.22 Adding Two Tables of Numbers In Example 9.19 we developed a C program to calculate the sum of the corresponding elements of two tables of integers. That program required three separate two-dimensional arrays, which were defined and processed in the conventional manner. Here is a variation of the program, in which each two-dimensional array is defined as an array of pointers to a set of one-dimensional integer arrays.

```
/* calculate the sum of the elements in two tables of integers */

/* each 2-dimensional array is processed as an array
   of pointers to a set of 1-dimensional integer arrays */

#include <stdio.h>
#include <stdlib.h>

#define MAXROWS 20

/* function prototypes */
void readinput (int *a[MAXROWS], int nrows, int ncols);
void computesums(int *a[MAXROWS], int *b[MAXROWS],
                 int *c[MAXROWS], int nrows, int ncols);
void writeoutput(int *c[MAXROWS], int nrows, int ncols);

main()
{
    int row, nrows, ncols;

    /* pointer definitions */
    int *a[MAXROWS], *b[MAXROWS], *c[MAXROWS];

    printf("How many rows? ");
    scanf("%d", &nrows);
    printf("How many columns? ");
    scanf("%d", &ncols);

    /* allocate initial memory */
    for (row = 0; row < nrows; ++row) {
        a[row] = (int *) malloc (ncols * sizeof(int));
        b[row] = (int *) malloc (ncols * sizeof(int));
        c[row] = (int *) malloc (ncols * sizeof(int));
    }

    printf("\n\nFirst table:\n");
    readinput(a, nrows, ncols);

    printf("\n\nSecond table:\n");
    readinput(b, nrows, ncols);

    computesums(a, b, c, nrows, ncols);

    printf("\n\nSums of the elements:\n\n");
    writeoutput(c, nrows, ncols);
}
```

```

void readinput(int *a[MAXROWS], int m, int n)
/* read in a table of integers */

{
    int row, col;

    for (row = 0; row < m; ++row) {
        printf("\nEnter data for row no. %2d\n", row + 1);
        for (col = 0; col < n; ++col)
            scanf("%d", (*(a + row) + col));
    }
    return;
}

void computesums(int *a[MAXROWS], int *b[MAXROWS], int *c[MAXROWS], int m, int n)
/* add the elements of two integer tables */

{
    int row, col;

    for (row = 0; row < m; ++row)
        for (col = 0; col < n; ++col)
            *((c + row) + col) = *((a + row) + col) + *((b + row) + col);
    return;
}

void writeoutput(int *a[MAXROWS], int m, int n)
/* write out a table of integers */

{
    int row, col;

    for (row = 0; row < m; ++row) {
        for (col = 0; col < n; ++col)
            printf("%4d", *((a + row) + col));
        printf("\n");
    }
    return;
}

```

In this program **a**, **b** and **c** are each defined as an array of pointers to integers. Each array has a maximum of **MAXROWS** elements. The function prototypes and the formal argument declarations within the subordinate functions also represent the arrays in this manner.

Since each element of **a**, **b** and **c** is a pointer, we must provide each pointer with enough memory for each row of integer quantities, using the **malloc** function as described in Sec. 10.5. These memory allocations appear in **main**, after the values for **nrows** and **ncols** have been entered. Consider the first memory allocation; i.e.,

```
a[row] = (int *) malloc(ncols * sizeof(int));
```

In this statement **a[0]** points to the the first row. Similarly, **a[1]** points to the second row, **a[2]** points to the third row, and so on. Thus, each array element points to a block of memory large enough to store one row of integer quantities (**ncols** integer quantities). Similar memory allocations are written for the other two arrays.

The individual array elements are processed by repeated use of the indirection operator. In **readinput**, for example, each array element is referenced as

```
scanf("%d", (*(a + row) + col));
```

Similarly, the addition of the array elements within `computesums` is written as

```
*(*(c + row) + col) = *(*(a + row) + col) + *(*(b + row) + col);
```

and the first `printf` statement within `writeoutput` is written as

```
printf("%4d", *(*(a + row) + col));
```

We could, of course, have used the more conventional notation within the functions. Thus, in `readinput` we could have written

```
scanf("%d", a[row][col]);
```

instead of

```
scanf("%d", (*(a + row) + col));
```

Similarly, in `computesums` we could have written

```
c[row][col] = a[row][col] + b[row][col];
```

instead of

```
*(*(c + row) + col) = *(*(a + row) + col) + *(*(b + row) + col);
```

and in `writeoutput` we could have written

```
printf("%4d", a[row][col]);
```

rather than

```
printf("%4d", *(*(a + row) + col));
```

This program will generate output identical to that shown in Example 9.19 when executed with the same input data.

10.8 ARRAYS OF POINTERS

A multidimensional array can be expressed in terms of an array of pointers rather than a pointer to a group of contiguous arrays. In such situations the newly defined array will have one less dimension than the original multidimensional array. Each pointer will indicate the beginning of a separate $(n - 1)$ -dimensional array.

In general terms, a two-dimensional array can be defined as a one-dimensional array of pointers by writing

```
data-type *array[ expression 1];
```

rather than the conventional array definition,

```
data-type array[ expression 1 ][ expression 2 ];
```

Similarly, an n -dimensional array can be defined as an $(n - 1)$ -dimensional array of pointers by writing

```
data-type *array[ expression 1 ][ expression 2 ] . . . [ expression n-1];
```

rather than

```
data-type array[ expression 1 ][ expression 2 ] . . . [ expression n];
```

In these declarations *data-type* refers to the data type of the original *n*-dimensional array, *array* is the array name, and *expression 1*, *expression 2*, . . . , *expression n* are positive-valued integer expressions that indicate the maximum number of elements associated with each subscript.

Notice that the array name and its preceding asterisk are *not* enclosed in parentheses in this type of declaration. (Compare carefully with the pointer declarations presented in the last section.) Thus, a right-to-left rule first associates the pairs of square brackets with *array*, defining the named object as an array. The preceding asterisk then establishes that the array will contain pointers.

Moreover, note that the *last* (i.e., the rightmost) expression is omitted when defining an array of pointers, whereas the *first* (i.e., the leftmost) expression is omitted when defining a pointer to a group of arrays. (Again, compare carefully with the declarations presented in the last section.) You should understand the distinction between this type of declaration and that presented in the last section.

When an *n*-dimensional array is expressed in this manner, an individual array element within the *n*-dimensional array can be accessed by a single use of the indirection operator. The following example illustrates how this is done.

EXAMPLE 10.23 Suppose *x* is a two-dimensional integer array having 10 rows and 20 columns, as in Example 10.20. We can define *x* as a one-dimensional array of pointers by writing

```
int *x[10];
```

Hence, *x[0]* points to the beginning of the first row, *x[1]* points to the beginning of the second row, and so on. Note that the number of elements within each row is not explicitly specified.

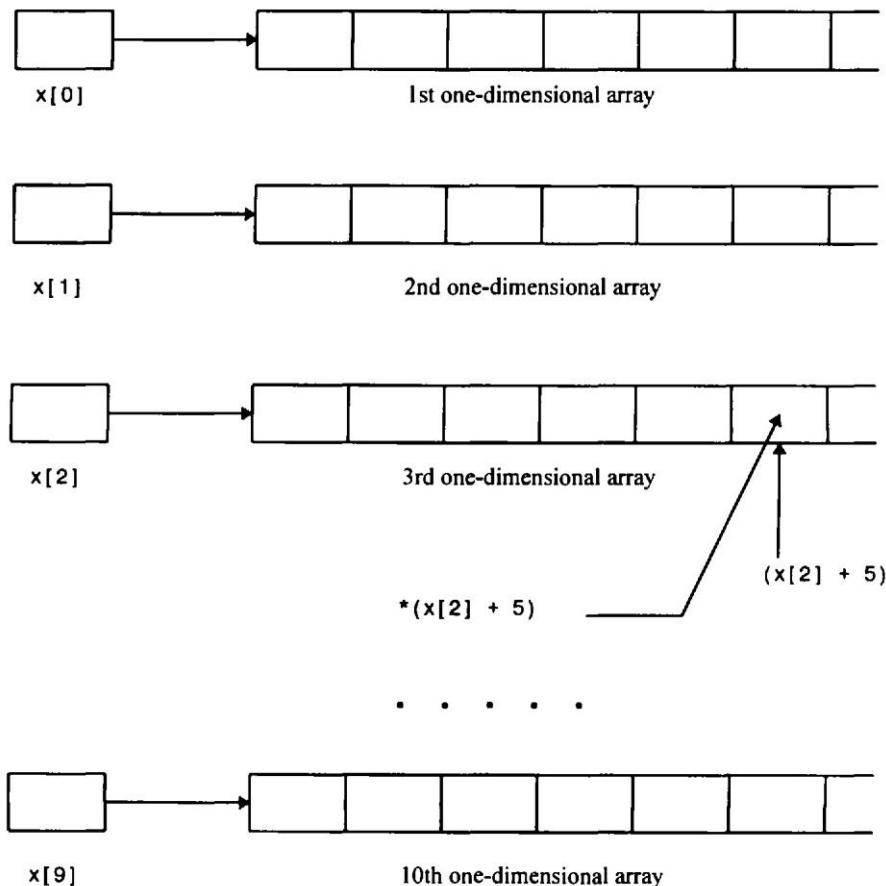


Fig. 10.6

An individual array element, such as $x[2][5]$, can be accessed by writing

```
*(x[2] + 5)
```

In this expression, $x[2]$ is a pointer to the first element in row 2, so that $(x[2] + 5)$ points to element 5 (actually, the sixth element) within row 2. The object of this pointer, $*(x[2] + 5)$, therefore refers to $x[2][5]$. These relationships are illustrated in Fig. 10.6.

Now consider a three-dimensional floating-point array t . Suppose the dimensionality of t is $10 \times 20 \times 30$. This array can be expressed as a two-dimensional array of pointers by writing

```
float *t[10][20];
```

Therefore we have 200 pointers (10 rows, 20 columns), each pointing to a one-dimensional array.

An individual array element, such as $t[2][3][5]$, can be accessed by writing

```
*(t[2][3] + 5)
```

In this expression, $t[2][3]$ is a pointer to the first element in the one-dimensional array represented by $t[2][3]$. Hence, $(t[2][3] + 5)$ points to element 5 (the sixth element) within this array. The object of this pointer, $*(t[2][3] + 5)$, therefore represents $t[2][3][5]$. This situation is directly analogous to the two-dimensional case described above.

EXAMPLE 10.24 Adding Two Tables of Numbers Here is yet another version of the programs presented in Examples 9.19 and 10.22, which calculate the sum of the corresponding elements of two tables of integers. Now each two-dimensional array is represented as an array of pointers to one-dimensional arrays. Each one-dimensional array will correspond to one row within the original two-dimensional array.

```
/* calculate the sum of the elements in two tables of integers */
/* each 2-dimensional array is represented as an array of pointers
   each pointer indicates a row in the original 2-dimensional array */

#include <stdio.h>
#include <stdlib.h>

#define MAXROWS 20
#define MAXCOLS 30

/* function prototypes */
void readinput(int *a[MAXROWS], int nrows, int ncols);
void computesums(int *a[MAXROWS], int *b[MAXROWS],
                 int *c[MAXROWS], int nrows, int ncols);
void writeoutput(int *c[MAXROWS], int nrows, int ncols);

main()
{
    int row, nrows, ncols;
    /* array definitions */
    int *a[MAXROWS], *b[MAXROWS], *c[MAXROWS];
    printf("How many rows? ");
    scanf("%d", &nrows);
    printf("How many columns? ");
    scanf("%d", &ncols);
```

```
/* allocate initial memory */
for (row = 0; row <= nrows; row++)  {
    a[row] = (int *) malloc(ncols * sizeof(int));
    b[row] = (int *) malloc(ncols * sizeof(int));
    c[row] = (int *) malloc(ncols * sizeof(int));
}

printf("\n\nFirst table:\n");
readinput(a, nrows, ncols);

printf("\n\nSecond table:\n");
readinput(b, nrows, ncols);

computesums(a, b, c, nrows, ncols);

printf("\n\nSums of the elements:\n\n");
writeoutput(c, nrows, ncols);
}

void readinput(int *a[MAXROWS], int m, int n)
/* read in a table of integers */
{
    int row, col;

    for (row = 0; row < m; ++row)  {
        printf('\nEnter data for row no. %2d\n', row + 1);
        for (col = 0; col < n; ++col)
            scanf("%d", (a[row] + col));
    }
    return;
}

void computesums(int *a[MAXROWS], int *b[MAXROWS],
                 int *c[MAXROWS], int m, int n)
/* add the elements of two integer tables */
{
    int row, col;

    for (row = 0; row < m; ++row)
        for (col = 0; col < n; ++col)
            *(c[row] + col) = *(a[row] + col) + *(b[row] + col);
    return;
}

void writeoutput(int *a[MAXROWS], int m, int n)
/* write out a table of integers */
{
    int row, col;

    for (row = 0; row < m; ++row)  {
        for (col = 0; col < n; ++col)
            printf("%4d", *(a[row] + col));
        printf("\n");
    }
    return;
}
```

Notice that `a`, `b` and `c` are now defined as one-dimensional arrays of pointers. Each array will contain `MAXROWS` elements (i.e., `MAXROWS` pointers). Each array element will point to a one-dimensional array of integers. The function prototypes and the formal argument declarations within the subordinate functions also represent the arrays in this manner.

Each one-dimensional array that is the object of a pointer (i.e., each row within each of the tables) must be allocated an initial block of memory. The `malloc` function accomplishes this. For example, each row within the first table is allocated an initial block of memory in the following manner.

```
a[row] = (int *) malloc(ncols * sizeof(int));
```

This statement associates a block of memory large enough to store `ncols` integer quantities with each pointer (i.e., with each element of `a`). Similar memory allocations are written for `b` and `c`. These `malloc` statements are placed within a `for` loop in order to allocate a block of memory for each of the nonzero rows within the three tables.

Notice the way the individual array elements are processed, using a combination of array and pointer notation. For example, in `readinput` each array element is now referenced as

```
scanf("%d", &(a[row] + col));
```

Similarly, in `computesums` and `writeoutput` the individual array elements are referenced as

```
*(c[row] + col) = *(a[row] + col) + *(b[row] + col);
```

and

```
printf("%4d", *(a[row] + col));
```

respectively. These statements could also have been written using conventional two-dimensional array notation.

This program, as well as the program presented in Example 10.22, will generate output that is identical to that shown in Example 9.19 when executed with the same input data. You may wish to verify this on your own. If this problem were being programmed from scratch, however, the conventional approach shown in Example 9.19, using two-dimensional arrays, would most likely be chosen.

Pointer arrays offer a particularly convenient method for storing strings. In this situation, each array element is a character-type pointer that indicates the beginning of a separate string. Thus, an n -element array can point to n different strings. Each individual string can be accessed by referring to its corresponding pointer.

EXAMPLE 10.25 Suppose the following strings are to be stored in a character-type array.

```
PACIFIC
ATLANTIC
INDIAN
CARIBBEAN
BERING
BLACK
RED
NORTH
BALTIc
CASPIAN
```

These strings can be stored in a two-dimensional character-type array; e.g.,

```
char names[10][12];
```

Note that `names` contains 10 rows, to accommodate the 10 strings. Each row must be large enough to store at least 10 characters, since `CARIBBEAN` contains 9 letters plus the null character (`\0`) at the end. To provide for larger strings, we are allowing each row to contain as many as 12 characters.

A better way to do this is to define a 10-element array of pointers; i.e.

```
char *names[10];
```

Thus, `names[0]` will point to `PACIFIC`, `names[1]` will point to `ATLANTIC`, and so on. Note that it is not necessary to include a maximum string size within the array declaration. However, a specified amount of memory will have to be allocated for each string later in the program, e.g.,

```
names[i] = (char *) malloc(12 * sizeof(char));
```

Just as individual strings can be accessed by referring to the corresponding pointer (i.e., the corresponding array element), so can individual string elements be accessed through the use of the indirection operator. For example, `*(*(names + 2) + 3)` refers to the fourth character (i.e., character number 3) in the third string (row number 2) of the array `names`, as defined in the preceding example.

Rearrangement of the strings can be accomplished simply by reassigning the pointers (i.e., by reassigning the elements in an array of pointers). The strings themselves need not be moved.

EXAMPLE 10.26 Reordering a List of Strings Consider once again the problem of entering a list of strings into the computer and rearranging them into alphabetical order. We saw one approach to this problem in Example 9.20, where the list of strings was stored in a two-dimensional array. Let us now approach this problem using a one-dimensional array of pointers, where each pointer indicates the beginning of a string. The string interchanges can now be carried out simply by reassigning the pointers, as required.

The complete program is presented below.

```
/* sort a list of strings into alphabetical order using an array of pointers */

#include <stdio.h>
#include <stdlib.h>
#include <string.h>

void reorder(int n, char *x[]);

main()
{
    int i, n = 0;
    char *x[10];

    printf("Enter each string on a separate line below\n\n");
    printf("Type '\\END\\' when finished\n\n");

    /* read in the list of strings */
    do {
        /* allocate memory */
        x[n] = (char *) malloc(12 * sizeof(char));
        printf("string %d: ", n + 1);
        scanf("%s", x[n]);
    }
    while (strcmp(x[n++], "END"));

    /* reorder the list of strings */
    reorder(--n, x);

    /* display the reordered list of strings */
    printf("\n\nReordered List of Strings:\n");
    for (i = 0; i < n; ++i)
        printf("\nstring %d: %s", i + 1, x[i]);
}
```

```

void reorder(int n, char *x[]) /* rearrange the list of strings */

{
    char *temp;
    int i, item;

    for (item = 0; item < n - 1; ++item)

        /* find the lowest of all remaining strings */
        for (i = item + 1; i < n; ++i)

            if (strcmp(x[item], x[i]) > 0) {
                /* interchange the two strings */
                temp = x[item];
                x[item] = x[i];
                x[i] = temp;
            }
    return;
}

```

The logic is essentially the same as that shown in Example 9.20, though the array containing the strings is now defined as an array of pointers. Notice that the second formal argument in the function `reorder` is declared in the same manner. Also, notice the string interchange routine (i.e., the `if` statement) within `reorder`. It is now the pointers, not the actual strings, that are interchanged. Hence the library function `strcpy`, which was used in Example 9.20, is not required. The program will therefore run somewhat faster than the earlier version.

Execution of this program will generate the same dialog as that shown in Example 9.20.

If the elements of an array are string pointers, a set of initial values can be specified as a part of the array declaration. In such cases the initial values will be strings, where each string corresponds to a separate array element. Remember, however, that an array must be declared `static` if it is initialized within a function.

An advantage to this scheme is that a fixed block of memory need not be reserved in advance, as is done when initializing a conventional array. Thus, if the initial declaration includes many strings and some of them are relatively short, there may be a substantial savings in memory allocation. Moreover, if some of the strings are particularly long, there is no need to worry about the possibility of exceeding some maximum specified string length (i.e., the maximum number of characters per row). Arrays of this type are often referred to as *ragged arrays*.

EXAMPLE 10.27 The following array declaration appears within a function.

```

static char *names[10] = {
    "PACIFIC",
    "ATLANTIC",
    "INDIAN",
    "CARIBBEAN",
    "BERING",
    "BLACK",
    "RED",
    "NORTH",
    "BALTIc",
    "CASPIAN"
};

```

In this example, `names` is a 10-element array of pointers. Thus, the first array element (i.e., the first pointer) will point to `PACIFIC`, the second array element will point to `ATLANTIC`, and so on.

Notice that the array is declared as **static** so that it can be initialized within the function. If the array declaration were external to all program functions, the **static** storage class designation would not be necessary.

Since the array declaration includes initial values, it is really not necessary to include an explicit size designation within the declaration. The size of the array will automatically be set equal to the number of strings that are present. Thus, the above declaration can be written as

```
static char *names[] = {
    "PACIFIC",
    "ATLANTIC",
    "INDIAN",
    "CARIBBEAN",
    "BERING",
    "BLACK",
    "RED",
    "NORTH",
    "BALTIC",
    "CASPIAN"
};
```

It should be understood that the ragged-array concept refers only to the *initialization* of string arrays, not the assignment of individual strings that may be read into the computer via the **scanf** function. Thus, applications requiring that strings be read into the computer and then processed, as in Example 10.26, still require the allocation of a specified amount of memory for each array element.

Initialized string values can be accessed by referring to their corresponding pointers (i.e., their corresponding array elements), in the usual manner. These pointers can be reassigned other string constants elsewhere in the program if necessary.

EXAMPLE 10.28 Displaying the Day of the Year Let us develop a program that will accept three integer quantities, indicating the month, day and year, and then display the corresponding day of the week, the month, the day and the year in a more legible manner. For example, suppose we were to enter the date 5 24 1997; this would produce the output Saturday, May 24, 1997. Programs of this type are often used to display information that is stored in a computer's internal memory in an encoded format.

Our basic strategy will be to enter a date into the computer, in the form *month, day, year* (**mm dd yyyy**), and then convert this date into the number of days relative to some base date. The day of the week corresponding to the specified date can then be determined quite easily, provided we know the day of the week corresponding to the base date. Let us arbitrarily choose Monday, January 1, 1900 as the base date. We will then convert any date beyond January 1, 1900 (actually, any date between January 1, 1900 and December 31, 2099) into an equivalent day of the week.

The computation can be carried out using the following empirical rules.

1. Enter numerical values for the variables **mm**, **dd** and **yy**, which represent the month, day and year, respectively (e.g., 5 24 1997).

2. Determine the approximate day of the current year, as

```
ndays = (long) (30.42 * (mm - 1)) + dd;
```

3. If **mm** == 2 (February), increase the value of **ndays** by 1.

4. If **mm** > 2 and **mm** < 8 (March, April, May, June or July), decrease the value of **ndays** by 1.

5. Convert the year into the number of years beyond the base date; i.e., **yy** -= 1900. Then test for a leap year as follows: If (**yy** % 4) == 0 and **mm** > 2, increase the value of **ndays** by 1.

6. Determine the number of complete 4-year cycles beyond the base date as **yy** / 4. For each complete 4-year cycle, add 1461 to **ndays**.

7. Determine the number of full years beyond the last complete 4-year cycle as **yy** % 4. For each full year, add 365 to **ndays**. Then add 1, because the first year beyond a full 4-year cycle will be a leap year.

8. If `ndays > 59` (i.e., if the date is any day beyond February 28, 1900), decrease the value of `ndays` by 1, because 1900 is not a leap year. (Note that the last year of each century is not a leap year, except those years that are evenly divisible by 400. Therefore 1900, the last year of the nineteenth century, is *not* a leap year, but 2000, the last year of the twentieth century, *is* a leap year.)
9. Determine the numerical day of the week corresponding to the specified date as `day = (ndays % 7)`.

Note that `day == 1` corresponds either to the base date, which is a Monday, or another date that also occurs on a Monday. Hence, `day == 2` will refer to a Tuesday, `day == 3` will refer to a Wednesday, . . . , `day == 6` will refer to a Saturday, and `day == 0` will refer to a Sunday.

Here is a complete function, called `convert`, that carries out steps 2 through 9. Note that `convert` accepts the integers `mm`, `dd` and `yy` as input parameters, and returns the integer quantity `(ndays % 7)`. Also, notice that `ndays` and `ncycles` are long integer variables, whereas all other variables are ordinary integers.

```
int convert(int mm, int dd, int yy) /* convert date to numerical day of week */

{
    long ndays;          /* number of days from start of 1900 */
    long ncycles;        /* number of 4-year cycles beyond 1900 */
    int nyears;          /* number of years beyond last 4-year cycle */
    int day;             /* day of week (0, 1, 2, 3, 4, 5 or 6) */

    /* numerical conversions */
    yy -= 1900;
    ndays = (long) (30.42 * (mm - 1)) + dd; /* approximate day of year */

    if (mm == 2) ++ndays;                      /* adjust for February */
    if ((mm > 2) && (mm < 8)) --ndays;        /* adjust for March - July */
    if ((yy % 4 == 0) && (mm > 2)) ++ndays;    /* adjust for leap year */

    ncycles = yy / 4;                         /* 4-year cycles beyond 1900 */
    ndays += ncycles * 1461;                   /* add days for 4-year cycles */

    nyyears = yy % 4;                         /* years beyond last 4-year cycle */
    if (nyyears > 0)                          /* add days for yrs beyond last 4-year cycle */
        ndays += 365 * nyyears + 1;

    if (ndays > 59) --ndays; /* adjust for 1900 (NOT a leap year) */
    day = ndays % 7;

    return(day);
}
```

The names of the days of the week can be stored as strings in a 7-element array; i.e.,

```
static char *weekday[] = {"Sunday", "Monday", "Tuesday", "Wednesday",
                         "Thursday", "Friday", "Saturday"};
```

Each day corresponds to the value assigned to `day`, where `day = (ndays % 7)`. The days begin with `Sunday` because `Sunday` corresponds to `day == 0`, as explained above. If the base date were not a Monday, this particular ordering of the days of the week would have to be changed.

Similarly, the names of the months can be stored as strings in a 12-element array; i.e.,

```
static char *month[] = {"January", "February", "March", "April", "May", "June", "July",
                       "August", "September", "October", "November", "December"};
```

Each month corresponds to the value of `mm - 1`.

Here is an entire C program that will carry out the conversion interactively.

```
/* convert a numerical date (mm dd yyyy) into 'day of week, month, day, year'
   (e.g., 5 24 1997 -> Saturday, May 24, 1997) */

#include <stdio.h>

void readininput(int *pm, int *pd, int *py);      /* function prototype */
int convert(int mm, int dd, int yy);                /* function prototype */

main()
{
    int mm, dd, yy;
    int day_of_week;      /* day of the week (0 -> Sunday,
                           1 -> Monday,
                           . . .
                           6 -> Saturday) */

    static char *weekday[] = {"Sunday", "Monday", "Tuesday", "Wednesday",
                             "Thursday", "Friday", "Saturday"};

    static char *month[] = {"January", "February", "March", "April",
                           "May", "June", "July", "August", "September",
                           "October", "November", "December"};

    /* opening message */
    printf("Date Conversion Routine\nTo STOP, enter 0 0 0");

    readininput(&mm, &dd, &yy);

    /* convert date to numerical day of week */
    while (mm > 0)
    {
        day_of_week = convert(mm, dd, yy);
        printf("\n%s %s %d, %d", weekday[day_of_week], month[mm-1], dd, yy);

        readininput(&mm, &dd, &yy);
    }
}

void readininput(int *pm, int *pd, int *py) /* read in the numerical date */
{
    printf("\n\nEnter mm dd yyyy: ");
    scanf("%d %d %d", pm, pd, py);
    return;
}

int convert(int mm, int dd, int yy) /* convert date to numerical day of week */
{
    long ndays;          /* number of days from start of 1900 */
    long ncycles;         /* number of 4-year cycles beyond 1900 */
    int nyears;           /* number of years beyond last 4-year cycle */
    int day;               /* day of week (0, 1, 2, 3, 4, 5 or 6) */
```

```

/* numerical conversions */
yy -= 1900;
ndays = (long) (30.42 * (mm - 1)) + dd; /* approximate day of year */

if (mm == 2) ++ndays; /* adjust for February */
if ((mm > 2) && (mm < 8)) --ndays; /* adjust for March - July */
if ((yy % 4 == 0) && (mm > 2)) ++ndays; /* adjust for leap year */

ncycles = yy / 4; /* 4-year cycles beyond 1900 */
ndays += ncycles * 1461; /* add days for 4-year cycles */

nyears = yy % 4; /* years beyond last 4-year cycle */
if (nyears > 0) /* add days for yrs beyond last 4-year cycle */
    ndays += 365 * nyears + 1;

if (ndays > 59) --ndays; /* adjust for 1900 (NOT a leap year) */
day = ndays % 7;

return(day);
}

```

This program includes a loop that repeatedly accepts a date in the form of three integers (i.e., mm dd yyyy) and returns the corresponding day and date in a more legible form. The program will continue to run until a value of 0 is entered for mm. Note that the prompt indicates that three zeros must be entered in order to stop the program execution; i.e., 0 0 0. Actually, however, the program only checks the value of mm.

A typical interactive session is shown below. As usual, the user's responses are underlined.

Date Conversion Routine

To STOP, enter 0 0 0

Enter mm dd yyyy: 10 29 1929

Tuesday, October 29, 1929

Enter mm dd yyyy: 8 15 1945

Wednesday, August 15, 1945

Enter mm dd yyyy: 7 20 1969

Sunday, July 20, 1969

Enter mm dd yyyy: 5 24 1997

Saturday, May 24, 1997

Enter mm dd yyyy: 8 30 2010

Monday, August 30, 2010

Enter mm dd yyyy: 4 12 2069

Friday, April 12, 2069

Enter mm dd yyyy: 0 0 0

10.9 PASSING FUNCTIONS TO OTHER FUNCTIONS

A pointer to a function can be passed to another function as an argument. This allows one function to be transferred to another, as though the first function were a variable. Let us refer to the first function as the *guest function*, and the second function as the *host function*. Thus, the guest is passed to the host, where it can be accessed. Successive calls to the host function can pass different pointers (i.e., different guest functions) to the host.

When a host function accepts the name of a guest function as an argument, the formal argument declaration must identify that argument as a pointer to the guest function. In its simplest form, the formal argument declaration can be written as

data-type (**function-name*)()

where *data-type* refers to the data type of the quantity returned by the guest and *function-name* is the name of the guest. The formal argument declaration can also be written as

data-type (**function-name*)(*type 1*, *type 2*, . . .)

or as

data-type (**function-name*)(*type 1 arg 1*, *type 2 arg 2*, . . .)

where *type 1*, *type 2*, . . . refer to the data types of the arguments associated with the guest, and *arg 1*, *arg 2*, . . . refer to the names of the arguments associated with the guest.

The guest function can be accessed within the host by means of the indirection operator. To do so, the indirection operator must precede the guest function name (i.e., the formal argument). Both the indirection operator and the guest function name must be enclosed in parentheses; i.e.,

(*function-name)(argument 1, argument 2, . . . , argument n);

where *argument 1*, *argument 2*, . . . , *argument n* refer to the arguments that are required in the function call.

Now consider the function declaration for the host function. It may be written as

funct-data-type *funct-name(arg-data-type) (*) (type₁, type₂, . . .)*,

← pointer to guest function →

data types of other funct args);

where *funct-data-type* refers to the data type of the quantity returned by the host function; *funct-name* refers to the name of the host function; *arg-data-type* refers to the data type of the quantity returned by the guest function, and *type 1*, *type 2*, . . . refer to the data types of guest function's arguments. Notice that the indirection operator appears in parentheses, to indicate a pointer to the guest function. Moreover, the data types of the guest function's arguments follow in a separate pair of parentheses, to indicate that they are function arguments.

When full function prototyping is used, the host function declaration is expanded as follows.

funct-data-type *funct-name*

(arg-data-type (*pt-var)(type 1 arg 1, type 2 arg 2, . . .),

← pointer to guest function →|

data types and names of other funct args);

The notation is the same as above, except that *pt-var* refers to the pointer variable pointing to the guest function, and *type 1 arg 1*, *type 2 arg 2*, . . . refer to the data types and the corresponding names of the guest function's arguments.

EXAMPLE 10.29 The skeletal outline of a C program is shown below. This program consists of four functions: *main*, *process*, *funct1* and *funct2*. Note that *process* is a host function for *funct1* and *funct2*. Each of the three subordinate functions returns an integer quantity.

```
int process(int (*)(int, int)); /* function declaration (host) */
int funct1(int, int);           /* function declaration (guest) */
int funct2(int, int);           /* function declaration (guest) */

main()
{
    int i, j;
    . . .
    i = process(funct1); /* pass funct1 to process; return a value for i */
    . . .
    j = process(funct2); /* pass funct2 to process; return a value for j */
    . . .
}

process(int (*pf)(int, int))      /* host function definition */
{                                /* (formal argument is a pointer to a function) */
    int a, b, c;
    . . .
    c = (*pf)(a, b); /* access the function passed to this function;
                        return a value for c */
    . . .
    return(c);
}

funct1(a, b)                    /* guest function definition */
int a, b;
{
    int c;
    c = . . . /* use a and b to evaluate c */
    return(c);
}

funct2(x, y)                    /* guest function definition */
int x, y;
{
    int z;
    z = . . . /* use x and y to evaluate z */
    return(z);
}
```

Notice that this program contains three function declarations. The declarations for `funct1` and `funct2` are straightforward. However the declaration for `process` requires some explanation. This declaration states that `process` is a host function that returns an integer quantity and has one argument. The argument is a pointer to a guest function that returns an integer quantity and has two integer arguments. The argument designation for the guest function is written as

```
int (*)(int, int)
```

Notice the way the argument designation fits into the entire host function declaration; i.e.,

```
int process(int (*)(int, int));
```

Now consider the formal argument declaration that appears within `process`; i.e.,

```
int (*pf)(int, int);
```

This declaration states that `pf` is a pointer to a guest function. The guest function will return an integer quantity, and it requires two integer arguments.

Here is another version of this same outline, utilizing full function prototyping. The changes are shown in boldface.

```
int process(int (*pf)(int a, int b)); /* function prototype (host) */
int funct1(int a, int b);           /* function prototype (guest) */
int funct2(int a, int b);           /* function prototype (guest) */

main()
{
    int i, j;
    . . .
    i = process(funct1); /* pass funct1 to process; return a value for i */
    . . .
    j = process(funct2); /* pass funct2 to process; return a value for j */
    . . .

}

process(int (*pf)(int a, int b)) /* host function definition */
{
    int a, b, c;
    . . .
    c = (*pf)(a, b); /* access the function passed to this function;
                        return a value for c */
    . . .
    return(c);
}

funct1(int a, int b) /* guest function definition */
{
    int c;
    c = . . .          /* use a and b to evaluate c */
    return(c);
}
```

```

funct2(int x, int y)      /* guest function definition */
{
    int z;

    z = . . .                /* use x and y to evaluate z */

    return(z);
}

```

The function prototypes include argument names as well as argument data types. Moreover, the prototype for **process** now includes the name of the variable (**p1**) that points to the guest function. Notice that the declaration of the formal argument **pf** within **process** is consistent with the function prototype.

Some programming applications can be formulated quite naturally in terms of one function being passed to another. For example, one function might represent a mathematical equation, and the other might contain a computational strategy to process the equation. In such cases the function representing the equation might be passed to the function that processes the equation. This is particularly useful if the program contains several different mathematical equations, one of which is selected by the user each time the program is executed.

EXAMPLE 10.30 Future Value of Monthly Deposits (Compound Interest Calculations) Suppose a person decides to save a fixed amount of money at the end of every month for n years. If the money earns interest at i percent per year, then it is natural to ask how much money will accumulate after n years (i.e., after $12n$ monthly deposits). The answer, of course, depends upon how much money is deposited each month, the interest rate, and the frequency with which the interest is compounded. For example, if the interest is compounded annually, semiannually, quarterly or monthly, the future amount of money that will accumulate after n years is given by

$$F = \frac{12A}{m} \left[\frac{(1 + i/m)^{mn} - 1}{i/m} \right] = 12A \left[\frac{(1 + i/m)^{mn} - 1}{i} \right]$$

where F is the future accumulation, A is the amount of money deposited each month, i is the annual interest rate (expressed as a decimal), and m is the number of compounding periods per year (e.g., $m = 1$ for annual compounding, $m = 2$ for semiannual compounding, $m = 4$ for quarterly compounding and $m = 12$ for monthly compounding).

If the compounding periods are shorter than the payment periods, such as in the case of daily compounding, then the future amount of money is determined by

$$F = A \left[\frac{(1 + i/m)^{mn} - 1}{(1 + i/m)^{m/12} - 1} \right]$$

Note that m is customarily assigned a value of 360 when the interest is compounded daily.

Finally, in the case of continuous compounding, the future amount of money is determined as

$$F = A \left[\frac{e^{in} - 1}{e^{i/12} - 1} \right]$$

Suppose we wish to determine F as a function of the annual interest rate i , for given values of A , m and n . Let us develop a program that will read the required input data into **main**, and then carry out the calculations within a separate function, called **table**. Each of the three formulas for determining the ratio F/A will be placed in one of three independent functions, called **md1**, **md2** and **md3**, respectively. Thus, the program will consist of five different functions.

When **table** is called from **main**, one of the arguments passed to **table** will be the name of the function containing the appropriate formula, as indicated by an input parameter (**freq**). The values of A , m and n that are read into **main** will also be passed to **table** as arguments. A loop will then be initiated within **table**, in which values of F are determined

for interest rates ranging from 0.01 (i.e., 1 percent per year) to 0.20 (20 percent per year). The calculated values will be displayed as they are generated. The entire program is shown below.

```
/* personal finance calculations */

#include <stdio.h>
#include <stdlib.h>
#include <ctype.h>
#include <math.h>

/* function prototypes */
void table (double (*pf)(double i, int m, double n), double a, int m, double n);
double md1(double i, int m, double n);
double md2(double i, int m, double n);
double md3(double i, int m, double n);

main()      /* calculate the future value of a series of monthly deposits */
{
    int m;          /* number of compounding periods per year */
    double n;        /* number of years */
    double a;        /* amount of each monthly payment */
    char freq;       /* frequency of compounding indicator */

    /* enter input data */
    printf("\nFUTURE VALUE OF A SERIES OF MONTHLY DEPOSITS\n\n");
    printf("Amount of Each Monthly Payment: ");
    scanf("%lf", &a);
    printf("Number of Years: ");
    scanf("%lf", &n);

    /* enter frequency of compounding */
    do  {
        printf("Frequency of Compounding (A, S, Q, M, D, C): ");
        scanf("%1s", &freq);
        freq = toupper(freq);      /* convert to upper case */
        if (freq == 'A')    {
            m = 1;
            printf("\nAnnual Compounding\n");
        }
        else if (freq == 'S')   {
            m = 2;
            printf("\nSemiannual Compounding\n");
        }
        else if (freq == 'Q')   {
            m = 4;
            printf("\nQuarterly Compounding\n");
        }
        else if (freq == 'M')   {
            m = 12;
            printf("\nMonthly Compounding\n");
        }
        else if (freq == 'D')   {
            m = 360;
            printf("\nDaily Compounding\n");
        }
    }
```

```

        else if (freq == 'C')    {
            m = 0;
            printf("\nContinuous Compounding\n");
        }
        else
            printf("\nERROR - Please Repeat\n\n");
    } while (freq != 'A' && freq != 'S' && freq != 'Q' &&
           freq != 'M' && freq != 'D' && freq != 'C');

/* carry out the calculations */
if (freq == 'C')
    table(md3, a, m, n); /* continuous compounding */
else if (freq == 'D')
    table(md2, a, m, n); /* daily compounding */
else
    table(md1, a, m, n); /* annual, semiannual, quarterly or monthly compounding */
}

void table (double (*pf)(double i, int m, double n), double a, int m, double n)
/* table generator (this function accepts a pointer to another function as an argument)

NOTE: double (*pf)(double i, int m, double n) is a POINTER TO A FUNCTION */

{
    int count; /* loop counter */
    double i; /* annual interest rate */
    double f; /* future value */

    printf("\nInterest Rate      Future Amount\n\n");
    for (count = 1; count <= 20; ++count) {
        i = 0.01 * count;
        f = a * (*pf)(i, m, n); /* ACCESS THE FUNCTION PASSED AS A POINTER */
        printf("      %2d          %.2f\n", count, f);
    }
    return;
}

double md1(double i, int m, double n)
/* monthly deposits, periodic compounding */

{
    double factor, ratio;

    factor = 1 + i/m;
    ratio = 12 * (pow(factor, m*n) - 1) / i;
    return(ratio);
}

double md2(double i, int m, double n)
/* monthly deposits, daily compounding */

{
    double factor, ratio;

    factor = 1 + i/m;
    ratio = (pow(factor, m*n) - 1) / (pow(factor, m/12) - 1);
    return(ratio);
}

```

```

double md3(double i, int dummy, double n)
/* monthly deposits, continuous compounding */
{
    double ratio;

    ratio = (exp(i*n) - 1) / (exp(i/12) - 1);
    return(ratio);
}

```

Notice the function prototypes, particularly the prototype for `table`. The first argument passed to `table` is a pointer to a guest function that receives two double-precision arguments and an integer argument, and returns a double-precision quantity. This pointer is intended to represent `md1`, `md2` or `md3`. The prototypes for these three functions follow the prototype for `table`. Each of these functions accepts two double-precision arguments and an integer argument, and returns a double-precision quantity, as required.

An interactive dialog for the input data is generated within `main`. In particular, the program accepts numerical values for `a` and `n`. It also accepts a one-character string for the character variable `freq`, which indicates the frequency of compounding. The only allowable characters that can be assigned to `freq` are A, S, Q, M, D or C (for Annual, Semiannual, Quarterly, Monthly, Daily or Continuous compounding, respectively). This character can be entered in either upper- or lowercase, since it is converted to uppercase within the program. Note that the program contains an error trap preventing characters other than A, S, Q, M, D or C from being entered.

An appropriate numerical value is assigned to `m` as soon as the frequency of compounding is determined. The program then accesses `table`, passing either `md1`, `md2` or `md3` as an argument, as determined by the character assigned to `freq`. (See the multiple `if - else` statement at the end of `main`.)

Now examine the host function `table`. The last three formal arguments (`a`, `m` and `n`) are declared as ordinary double-precision or integer variables. However, the first formal argument (`pf`) is declared as a pointer to a guest function that accepts two double-precision arguments and an integer argument, and returns a double-precision quantity. These formal argument declarations are consistent with the function prototype for `table`.

The values for `i` (i.e., the interest rates) are generated internally within `table`. These values are determined as `0.01 * count`. Since `count` ranges from 1 to 20, we see that the interest rates range from 0.01 to 0.20, as required.

Notice the manner in which the values for `f` are calculated within `table`; i.e.,

```
f = a * (*pf)(i, m, n);
```

The expression `(*pf)` refers to the guest function whose name is passed to `table` (i.e., either `md1`, `md2` or `md3`). This is accompanied by the required list of arguments, containing the current values for `i`, `m` and `n`. The value returned by the guest function is then multiplied by `a`, and the product is assigned to `f`.

The three remaining functions, `md1`, `md2` and `md3`, are straightforward. Notice that the second argument in `md3` is called `dummy`, because the value of this argument is not utilized within the function. We could have done this with `md2` as well, since the value of `m` is always 360 in the case of daily compounding.

Execution of the program produces the following representative dialog.

FUTURE VALUE OF A SERIES OF MONTHLY DEPOSITS

Amount of Each Monthly Payment: 100

Number of Years: 3

Frequency of Compounding (A, S, Q, M, D, C): D

ERROR - Please Repeat

Frequency of Compounding (A, S, Q, M, D, C): M

Monthly Compounding

Interest Rate	Future Amount
1	3653.00
2	3707.01
3	3762.06
4	3818.16
5	3875.33
6	3933.61
7	3993.01
8	4053.56
9	4115.27
10	4178.18
11	4242.31
12	4307.69
13	4374.33
14	4442.28
15	4511.55
16	4582.17
17	4654.18
18	4727.60
19	4802.45
20	4878.78

10.10 MORE ABOUT POINTER DECLARATIONS

Before leaving this chapter we mention that pointer declarations can become complicated, and some care is required in their interpretation. This is especially true of declarations that involve functions or arrays.

One difficulty is the dual use of parentheses. In particular, parentheses are used to indicate functions, and they are used for nesting purposes (to establish precedence) within more complicated declarations. Thus, the declaration

```
int *p(int a);
```

indicates a function that accepts an integer argument, and returns a pointer to an integer. On the other hand, the declaration

```
int (*p)(int a);
```

indicates a *pointer to a function* that accepts an integer argument and returns an integer. In this last declaration, the first pair of parentheses is used for nesting, and the second pair is used to indicate a function.

The interpretation of more complex declarations can be increasingly troublesome. For example, consider the declaration

```
int *(*p)(int (*a)[]);
```

In this declaration, $(\ast p)(\dots)$ indicates a pointer to a function. Hence, $\text{int } \ast(\ast p)(\dots)$ indicates a pointer to a function that returns a pointer to an integer. Within the last pair of parentheses (the function's argument specification), $(\ast a)[]$ indicates a pointer to an array. Therefore, $\text{int } (\ast a)[]$ represents a pointer to an array of integers. Putting the pieces together, $(\ast p)(\text{int } (\ast a)[])$ represents a pointer to a function whose argument is a pointer to an array of integers. And finally, the entire declaration

```
int *(*p)(int (*a)[]);
```

represents a pointer to a function that accepts a pointer to an array of integers as an argument, and returns a pointer to an integer.

Remember that a left parenthesis immediately following an identifier name indicates that the identifier represents a function. Similarly, a left square bracket immediately following an identifier name indicates that the identifier represents an array. Parentheses that identify functions and square brackets that identify arrays have a higher precedence than the unary indirection operator (see Appendix C). Therefore, additional parentheses are required when declaring a pointer to a function or a pointer to an array.

The following example provides a number of illustrations.

EXAMPLE 10.31 Several declarations involving pointers are shown below. The individual declarations range from simple to complex.

```
int *p;                      /* p is a pointer to an integer quantity */

int *p[10];                  /* p is a 10-element array of pointers to integer quantities */

int (*p)[10];                /* p is a pointer to a 10-element integer array */

int *p(void);                /* p is a function that
                                returns a pointer to an integer quantity */

int p(char *a);              /* p is a function that
                                accepts an argument which is a pointer to a character and
                                returns an integer quantity */

int *p(char a*);             /* p is a function that
                                accepts an argument which is a pointer to a character
                                returns a pointer to an integer quantity */

int (*p)(char *a);           /* p is a pointer to a function that
                                accepts an argument which is a pointer to a character
                                returns an integer quantity */

int (*p(char *a))[10];        /* p is a function that
                                accepts an argument which is a pointer to a character
                                returns a pointer to a 10-element integer array */

int p(char (*a)[]);           /* p is a function that
                                accepts an argument which is a pointer to a character array
                                returns an integer quantity */

int p(char *a[]);             /* p is a function that
                                accepts an argument which is an array of pointers to
                                characters
                                returns an integer quantity */

int *p(char a[]);             /* p is a function that
                                accepts an argument which is a character array
                                returns a pointer to an integer quantity */

int *p(char (*a)[]);          /* p is a function that
                                accepts an argument which is a pointer to a character array
                                returns a pointer to an integer quantity */

int *p(char *a[]);             /* p is a function that
                                accepts an argument which is a character array
                                returns a pointer to an integer quantity */
```

```
accepts an argument which is an array of pointers to
characters
returns a pointer to an integer quantity */

int (*p)(char (*a)[]); /* p is a pointer to a function that
                           accepts an argument which is a pointer to a character array
                           returns an integer quantity */

int *(*p)(char (*a)[]); /* p is pointer to a function that
                           accepts an argument which is a pointer to a character array
                           returns a pointer to an integer quantity */

int *(*p)(char *a[]); /* p is a pointer to a function that
                           accepts an argument which is an array of pointers to
                           characters
                           returns a pointer to an integer quantity */

int (*p[10])(void); /* p is a 10-element array of pointers to functions;
                           each function returns an integer quantity */

int (*p[10])(char a); /* p is a 10-element array of pointers to functions;
                           each function accepts an argument which is a character, and
                           returns an integer quantity */

int *(*p[10])(char a); /* p is a 10-element array of pointers to functions;
                           each function accepts an argument which is a character, and
                           returns a pointer to an integer quantity */

int *(*p[10])(char *a); /* p is a 10-element array of pointers to functions;
                           each function accepts an argument which is a pointer to a
                           character, and
                           returns a pointer to an integer quantity */
```

Review Questions

- 10.1 For the version of C available on your particular computer, how many memory cells are required to store a single character? An integer quantity? A long integer? A floating-point quantity? A double-precision quantity?
- 10.2 What is meant by the address of a memory cell? How are addresses usually numbered?
- 10.3 How is a variable's address determined?
- 10.4 What kind of information is represented by a pointer variable?
- 10.5 What is the relationship between the address of a variable v and the corresponding pointer variable pv?
- 10.6 What is the purpose of the indirection operator? To what type of operand must the indirection operator be applied?
- 10.7 What is the relationship between the data item represented by a variable v and the corresponding pointer variable pv?
- 10.8 What precedence is assigned to the unary operators compared with the multiplication, division and remainder operators? In what order are the unary operators evaluated?
- 10.9 Can the address operator act upon an arithmetic expression, such as $2 * (u + v)$? Explain the reasons for your answer.
- 10.10 Can an expression involving the indirection operator appear on the left side of an assignment statement? Explain.

- 10.11 What kinds of objects can be associated with pointer variables?
- 10.12 How is a pointer variable declared? What is the purpose of the data type included in the declaration?
- 10.13 In what way can the assignment of an initial value be included in the declaration of a pointer variable?
- 10.14 Are integer values ever assigned to pointer variables? Explain.
- 10.15 Why is it sometimes desirable to pass a pointer to a function as an argument?
- 10.16 Suppose a function receives a pointer as an argument. Explain how the function prototype is written. In particular, explain how the data type of the pointer argument is represented.
- 10.17 Suppose a function receives a pointer as an argument. Explain how the pointer argument is declared within the function definition.
- 10.18 What is the relationship between an array name and a pointer? How is an array name interpreted when it appears as an argument to a function?
- 10.19 Suppose a formal argument within a function definition is an array. How can the array be declared within the function?
- 10.20 How can a portion of an array be passed to a function?
- 10.21 How can a function return a pointer to its calling routine?
- 10.22 Describe two different ways to specify the address of an array element.
- 10.23 Why is the value of an array subscript sometimes referred to as an offset when the subscript is a part of an expression indicating the address of an array element?
- 10.24 Describe two different ways to access an array element. Compare your answer to that of question 10.22.
- 10.25 Can an address be assigned to an array name or an array element? Can an address be assigned to a pointer variable whose object is an array?
- 10.26 Suppose a numerical array is defined in terms of a pointer variable. Can the individual array elements be initialized?
- 10.27 Suppose a character-type array is defined in terms of a pointer variable. Can the individual array elements be initialized? Compare your answer with that of the previous question.
- 10.28 What is meant by dynamic memory allocation? What library function is used to allocate memory dynamically? How is the size of the memory block specified? What kind of information is returned by the library function?
- 10.29 Suppose an integer quantity is added to or subtracted from a pointer variable. How will the sum or difference be interpreted?
- 10.30 Under what conditions can one pointer variable be subtracted from another? How will this difference be interpreted?
- 10.31 Under what conditions can two pointer variables be compared? Under what conditions are such comparisons useful?
- 10.32 How is a multidimensional array defined in terms of a pointer to a collection of contiguous arrays of lower dimensionality?
- 10.33 How can the indirection operator be used to access a multidimensional array element?
- 10.34 How is a multidimensional array defined in terms of an array of pointers? What does each pointer represent? How does this definition differ from a pointer to a collection of contiguous arrays of lower dimensionality?
- 10.35 How can a one-dimensional array of pointers be used to represent a collection of strings?
- 10.36 If several strings are stored within a one-dimensional array of pointers, how can an individual string be accessed?
- 10.37 If several strings are stored within a one-dimensional array of pointers, what happens if the strings are reordered? Are the strings actually moved to different locations within the array?
- 10.38 Under what conditions can the elements of a multidimensional array be initialized if the array is defined in terms of an array of pointers?
- 10.39 When transferring one function to another, what is meant by the guest function? What is the host function?

- 10.40 Suppose a formal argument within a host function definition is a pointer to another function. How is the formal argument declared? Within the declaration, to what does the data type refer?
- 10.41 Suppose a formal argument within the definition of the host function p is a pointer to the guest function q. How is the formal argument declared within p? In this declaration, to what does the data type refer? How is function q accessed within function p?
- 10.42 Suppose that p is a host function, and one of p's arguments is a pointer to function q. How would the declaration for p be written if full function prototyping is used?
- 10.43 For what types of applications is it particularly useful to pass one function to another?

Problems

- 10.44 Explain the meaning of each of the following declarations.

- (a) int *px;
- (b) float a, b;
float *pa, *pb;
- (c) float a = -0.167;
float *pa = &a;
- (d) char c1, c2, c3;
char *pc1, *pc2, *pc3 = &c1;
- (e) double funct(double *a, double *b, int *c);
- (f) double *funct(double *a, double *b, int *c);
- (g) double (*a)[12];
- (h) double *a[12];
- (i) char *a[12];
- (j) char *d[4] = {"north", "south", "east", "west"};
- (k) long (*p)[10][20];
- (l) long *p[10][20];
- (m) char sample(int (*pf)(char a, char b));
- (n) int (*pf)(void);
- (o) int (*pf)(char a, char b);
- (p) int (*pf)(char *a, char *b);

- 10.45 Write an appropriate declaration for each of the following situations.

- (a) Declare two pointers whose objects are the integer variables i and j.
- (b) Declare a pointer to a floating-point quantity, and a pointer to a double-precision quantity.
- (c) Declare a function that accepts two integer arguments and returns a pointer to a long integer.
- (d) Declare a function that accepts two arguments and returns a long integer. Each argument will be a pointer to an integer quantity.
- (e) Declare a one-dimensional floating-point array using pointer notation.
- (f) Declare a two-dimensional floating-point array, with 15 rows and 30 columns, using pointer notation.
- (g) Declare an array of strings whose initial values are “red,” “green” and “blue.”
- (h) Declare a function that accepts another function as an argument and returns a pointer to a character. The function passed as an argument will accept an integer argument and return an integer quantity.
- (i) Declare a pointer to a function that accepts three integer arguments and returns a floating-point quantity.
- (j) Declare a pointer to a function that accepts three pointers to integer quantities as arguments and returns a pointer to a floating-point quantity.

10.46 A C program contains the following statements.

```
char u, v = 'A';
char *pu, *pv = &v;

. . .

*pv = v + 1;
u = *pv + 1;
pu = &u;
```

Suppose each character occupies 1 byte of memory. If the value assigned to *u* is stored in (hexadecimal) address F8C and the value assigned to *v* is stored in address F8D, then

- (a) What value is represented by *&v*?
- (b) What value is assigned to *pv*?
- (c) What value is represented by **pv*?
- (d) What value is assigned to *u*?
- (e) What value is represented by *&u*?
- (f) What value is assigned to *pu*?
- (g) What value is represented by **pu*?

10.47 A C program contains the following statements.

```
int i, j = 25;
int *pi, *pj = &j;

. . .

*pj = j + 5;
i = *pj + 5;
pi = pj;
*pi = i + j;
```

Suppose each integer quantity occupies 2 bytes of memory. If the value assigned to *i* begins at (hexadecimal) address F9C and the value assigned to *j* begins at address F9E, then

- (a) What value is represented by *&i*?
- (b) What value is represented by *&j*?
- (c) What value is assigned to *pj*?
- (d) What value is assigned to **pj*?
- (e) What value is assigned to *i*?
- (f) What value is represented by *pi*?
- (g) What final value is assigned to **pi*?
- (h) What value is represented by *(pi + 2)*?
- (i) What value is represented by the expression *(*pi + 2)*?
- (j) What value is represented by the expression **(pi + 2)*?

10.48 A C program contains the following statements.

```
float a = 0.001, b = 0.003;
float c, *pa, *pb;

pa = &a;
*pa = 2 * a;
pb = &b;
c = 3 * (*pb - *pa);
```

Suppose each floating-point number occupies 4 bytes of memory. If the value assigned to **a** begins at (hexadecimal) address 1130, the value assigned to **b** begins at address 1134, and the value assigned to **c** begins at 1138, then

- (a) What value is assigned to **&a**?
- (b) What value is assigned to **&b**?
- (c) What value is assigned to **&c**?
- (d) What value is assigned to **pa**?
- (e) What value is represented by ***pa**?
- (f) What value is represented by **&(*pa)**?
- (g) What value is assigned to **pb**?
- (h) What value is represented by ***pb**?
- (i) What value is assigned to **c**?

10.49 The skeletal structure of a C program is shown below.

```

int funct1(char a, char b);
int funct2(char *pa, char *pb);

main()
{
    char a = 'X';
    char b = 'Y';
    int i, j;

    . . . .

    i = funct1(a, b);
    printf("a=%c    b=%c\n", a, b);

    . . . .

    j = funct2(&a, &b);
    printf("a=%c    b=%c", a, b);
}

int funct1(char c1, char c2)
{
    c1 = 'P';
    c2 = 'Q';

    . . . .

    return((c1 < c2) ? c1 : c2);
}

int funct2(char *c1, char *c2)
{
    *c1 = 'P';
    *c2 = 'Q';

    . . . .

    return((*c1 == *c2) ? *c1 : *c2);
}

```

- (a) Within main, what value is assigned to **i**?
- (b) What value is assigned to **j**?

- (c) What values are displayed by the first `printf` statement?
- (d) What values are displayed by the second `printf` statement?

Assume ASCII characters.

10.50 The skeletal structure of a C program is shown below.

```
void funct(int *p);

main()
{
    static int a[5] = {10, 20, 30, 40, 50};

    . . . .

    funct(a);

    . . . .

}

void funct(int *p)
{
    int i, sum = 0;
    for (i = 0; i < 5; ++i)
        sum += *(p + i);
    printf("sum=%d", sum);
    return;
}
```

- (a) What kind of argument is passed to `funct`?
- (b) What kind of information is returned by `funct`?
- (c) What kind of formal argument is defined within `funct`?
- (d) What is the purpose of the `for` loop that appears within `funct`?
- (e) What value is displayed by the `printf` statement within `funct`?

10.51 The skeletal structure of a C program is shown below.

```
void funct(int *p);

main()
{
    static int a[5] = {10, 20, 30, 40, 50};

    . . . .

    funct(a + 3);

    . . . .

}

void funct(int *p)
{
    int i, sum = 0;
    for (i = 3; i < 5; ++i)
        sum += *(p + i);
    printf("sum=%d", sum);
    return;
}
```

- (a) What kind of argument is passed to **funct**?
- (b) What kind of information is returned by **funct**?
- (c) What information is actually passed to **funct**?
- (d) What is the purpose of the **for** loop that appears within **funct**?
- (e) What value is displayed by the **printf** statement within **funct**?

Compare your answers with those of the previous problem. In what ways do these two skeletal outlines differ?

- 10.52** The skeletal structure of a C program is shown below.

```

int *funct(int *p);

main()
{
    static int a[5] = {10, 20, 30, 40, 50};
    int *ptmax;

    . . . .

    ptmax = funct(a);
    printf("max=%d", *ptmax);

    . . . .

}

int *funct(int *p)
{
    int i, imax, max = 0;
    for (i = 0; i < 5; ++i)
        if (*(p + i) > max) {
            max = *(p + i);
            imax = i;
        }
    return(p + imax);
}

```

- (a) Within **main**, what is **ptmax**?
- (b) What kind of information is returned by **funct**?
- (c) What is assigned to **ptmax** when the function is accessed?
- (d) What is the purpose of the **for** loop that appears within **funct**?
- (e) What value is displayed by the **printf** statement within **main**?

Compare your answers with those of the previous two problems. In what ways are the skeletal outlines different?

- 10.53** A C program contains the following declaration.

```
static int x[8] = {10, 20, 30, 40, 50, 60, 70, 80};
```

- (a) What is the meaning of **x**?
- (b) What is the meaning of **(x + 2)**?
- (c) What is the value of ***x**?
- (d) What is the value of **(*x + 2)**?
- (e) What is the value of ***(x + 2)**?

10.54 A C program contains the following declaration.

```
static float table[2][3] = {  
    {1.1, 1.2, 1.3},  
    {2.1, 2.2, 2.3}  
};
```

- (a) What is the meaning of `table`?
- (b) What is the meaning of `(table + 1)`?
- (c) What is the meaning of `*(table + 1)`?
- (d) What is the meaning of `(*table + 1) + 1`?
- (e) What is the meaning of `(*table) + 1`?
- (f) What is the value of `*(*table + 1) + 1`?
- (g) What is the value of `*(*table) + 1`?
- (h) What is the value of `*(*table + 1))`?
- (i) What is the value of `*(*table) + 1` + 1?

10.55 A C program contains the following declaration.

```
static char *color[6] = {"red", "green", "blue", "white", "black", "yellow"};
```

- (a) What is the meaning of `color`?
- (b) What is the meaning of `(color + 2)`?
- (c) What is the value of `*color`?
- (d) What is the value of `*(color + 2)`?
- (e) How do `color[5]` and `*(color + 5)` differ?

10.56 The skeletal structure of a C program is shown below.

```
float one(float x, float y);  
float two(float x, float y);  
float three(float (*pt)(float x, float y));  
  
main()  
{  
    float a, b;  
    . . . . .  
    a = three(one);  
    . . . . .  
    b = three(two);  
    . . . . .  
}  
  
float one(float x, float y)  
{  
    float z;  
    z = . . . . .  
    return(z);  
}
```

```
float two(float p, float q)
{
    float r;
    r = . . . . .
    return(r);
}

float three(float (*pt)(float x, float y))
{
    float a, b, c;
    . . . . .
    c = (*pt)(a, b);
    . . . . .
    return(c);
}
```

- (a) Interpret each of the function prototypes.
- (b) Interpret the definitions of the functions one and two.
- (c) Interpret the definition of the function three. How does three differ from one and two?
- (d) What happens within main each time three is accessed?

10.57 The skeletal structure of a C program is shown below.

```
float one(float *px, float *py);
float two(float *px, float *py);
float *three(float (*pt)(float *px, float *py));

main()
{
    float *pa, *pb;
    . . . . .
    pa = three(one);
    . . . . .
    pb = three(two);
    . . . . .
}

float one(float *px, float *py)
{
    float z;
    z = . . . . .
    return(z);
}
```

```
float two(float *pp, float *pq)
{
    float r;
    r = . . . . .
    return(r);
}

float *three(float (*pt)(float *px, float *py))
{
    float a, b, c;
    . . . . .
    c = (*pt)(&a, &b);
    . . . . .
    return(&c);
}
```

- (a) Interpret each of the function prototypes.
(b) Interpret the definitions of the functions `one` and `two`.
(c) Interpret the definition of the function `three`. How does `three` differ from `one` and `two`?
(d) What happens within `main` each time `three` is accessed?
(e) How does this program outline differ from the outline shown in the last example?

10.58 Explain the purpose of each of the following declarations.

- (a) `float (*x)(int *a);`
- (b) `float (*x(int *a))[20];`
- (c) `float x(int (*a)[]);`
- (d) `float x(int *a[]);`
- (e) `float *x(int a[]);`
- (f) `float *x(int (*a)[]);`
- (g) `float *x(int *a[]);`
- (h) `float (*x)(int (*a)[]);`
- (i) `float *(*x)(int *a[]);`
- (j) `float (*x[20])(int a);`
- (k) `float *(*x[20])(int *a);`

10.59 Write an appropriate declaration for each of the following situations involving pointers.

- (a) Declare a function that accepts an argument which is a pointer to an integer quantity and returns a pointer to a six-element character array.
- (b) Declare a function that accepts an argument which is a pointer to an integer array and returns a character.
- (c) Declare a function that accepts an argument which is an array of pointers to integer quantities and returns a character.
- (d) Declare a function that accepts an argument which is an integer array and returns a pointer to a character.
- (e) Declare a function that accepts an argument which is a pointer to an integer array and returns a pointer to a character.
- (f) Declare a function that accepts an argument which is an array of pointers to integer quantities and returns a pointer to a character.

- (g) Declare a pointer to a function that accepts an argument which is a pointer to an integer array and returns a character.
- (h) Declare a pointer to a function that accepts an argument which is a pointer to an integer array and returns a pointer to a character.
- (i) Declare a pointer to a function that accepts an argument which is an array of pointers to integer quantities and returns a pointer to a character.
- (j) Declare a 12-element array of pointers to functions. Each function will accept two double-precision quantities as arguments and will return a double-precision quantity.
- (k) Declare a 12-element array of pointers to functions. Each function will accept two double-precision quantities as arguments and will return a pointer to a double-precision quantity.
- (l) Declare a 12-element array of pointers to functions. Each function will accept two pointers to double-precision quantities as arguments and will return a pointer to a double-precision quantity.

Programming Problems

10.60 Modify the program shown in Example 10.1 as follows.

- (a) Use floating-point data rather than integer data. Assign an initial value of 0.3 to u.
- (b) Use double-precision data rather than integer data. Assign an initial value of 0.3×10^{45} to u.
- (c) Use character data rather than integer data. Assign an initial value of 'C' to u.

Execute each modification and compare the results with those given in Example 10.1. Be sure to modify the `printf` statements accordingly.

10.61 Modify the program shown in Example 10.3 as follows.

- (a) Use floating-point data rather than integer data. Assign an initial value of 0.3 to v.
- (b) Use double-precision data rather than integer data. Assign an initial value of 0.3×10^{45} to v.
- (c) Use character data rather than integer data. Assign an initial value of 'C' to v.

Execute each modification and compare the results with those given in Example 10.3. Be sure to modify the `printf` statements accordingly.

10.62 Modify the program shown in Example 10.7 so that a single one-dimensional, character-type array is passed to `funct1`. Delete `funct2` and all references to `funct2`. Initially, assign the string "red" to the array within `main`. Then reassign the string "green" to the array within `funct1`. Execute the program and compare the results with those shown in Example 10.7. Remember to modify the `printf` statements accordingly.

10.63 Modify the program shown in Example 10.8 (analyzing a line of text) so that it also counts the number of words and the total number of characters in the line of text. (*Note:* A new word can be recognized by the presence of a blank space followed by a nonwhitespace character.) Test the program using the text given in Example 10.8.

10.64 Modify the program shown in Example 10.8 (analyzing a line of text) so that it can process multiple lines of text. First enter and store all lines of text. Then determine the number of vowels, consonants, digits, whitespace characters and "other" characters for each line. Finally, determine the average number of vowels per line, consonants per line, etc. Write and execute the program two different ways.

- (a) Store the multiple lines of text in a two-dimensional array of characters.
- (b) Store the multiple lines of text as individual strings whose maximum length is unspecified. Maintain a pointer to each string within a one-dimensional array of pointers.

In each case, identify the last line of text in some predetermined manner (e.g., by entering the string "END"). Test the program using several lines of text of your own choosing.

10.65 Modify the program shown in Example 10.12 so that the elements of x are long integers rather than ordinary integers. Execute the program and compare the results with those shown in Example 10.12. (Remember to modify the `printf` statement to accommodate the long integer quantities.)

- 10.66** Modify the program shown in Example 10.16 so that any one of the following rearrangements can be carried out:
- Smallest to largest, by magnitude
 - Smallest to largest, algebraic
 - Largest to smallest, by magnitude
 - Largest to smallest, algebraic

Use pointer notation to represent individual integer quantities, as in Example 10.16. (Recall that an array version of this problem was presented in Example 9.13.) Include a menu that will allow the user to select which rearrangement will be used each time the program is executed. Test the program using the following 10 values.

4.7	-8.0
-2.3	11.4
12.9	5.1
8.8	-0.2
6.0	-14.7

- 10.67** Modify the program shown in Example 10.22 (adding two tables of numbers) so that each element in the table *c* is the larger of the corresponding elements in tables *a* and *b* (rather than the sum of the corresponding elements in *a* and *b*). Represent each table (each array) as a pointer to a group of one-dimensional arrays, as in Example 10.22. Use pointer notation to access the individual table elements. Test the program using the tabular data provided in Example 9.19. (You may wish to experiment with this program, using several different ways to represent the arrays and the individual array elements.)
- 10.68** Repeat the previous problem, representing each table (each array) as a one-dimensional array of pointers, as discussed in Example 10.24.
- 10.69** Modify the program shown in Example 10.26 (reordering a list of strings) so that the list of strings can be rearranged into either alphabetical or reverse-alphabetical order. Use pointer notation to represent the beginning of each string. Include a menu that will allow the user to select which rearrangement will be used each time the program is executed. Test the program using the data provided in Example 9.20.
- 10.70** Modify the program shown in Example 10.28 (displaying the day of the year) so that it can determine the number of days between two dates, assuming both dates are beyond the base date of January 1, 1900. (*Hint:* Determine the number of days between the first specified date and the base date; then determine the number of days between the second specified date and the base date. Finally, determine the difference between these two calculated values.)
- 10.71** Modify the program shown in Example 10.30 (compound interest calculations) so that it generates a table of *F*-values for various interest rates, using different compounding frequencies. Assume that *A* and *n* are input values. Display the output in the following manner.

A = . . .
n = . . .
Interest rate = 5% 6% 7% 8% 9% 10% 11% 12% 13% 14% 15%
Frequency of Compounding
Annual
Semiannual
Quarterly
Monthly
Daily
Continuously

Notice that the first four rows are generated by one function with different arguments, and each of the last two rows is generated by a different function.

- 10.72** Modify the program shown in Example 10.30 (compound interest calculations) so that it generates a table of F -values for various time periods, using different compounding frequencies. Assume that A and i are input values. Display the output in the following manner.

	A = . . .
	i = . . .
Time period (n) =	1 2 3 4 5 6 7 8 9 10
Frequency of Compounding	
Annual	— — — — — — — — — —
Semiannual	— — — — — — — — — —
Quarterly	— — — — — — — — — —
Monthly	— — — — — — — — — —
Daily	— — — — — — — — — —
Continuously	— — — — — — — — — —

Notice that the first four rows are generated by one function with different arguments, and each of the last two rows is generated by a different function.

- 10.73** Repeat the previous problem, but transpose the table so that each row represents a different value for n and each column represents a different compounding frequency. Consider integer values of n ranging from 1 to 50. Note that this table will consist of 50 rows and 6 columns. (*Hint:* Generate the table by columns, storing each column in a two-dimensional array. Display the entire array after all the values have been generated.)

Compare the programming effort required for this problem with the programming effort required for the preceding problem.

- 10.74** Examples 9.8 and 9.9 present programs to calculate the average of a list of numbers and then calculate the deviations about the average. Both programs make use of one-dimensional, floating-point arrays. Modify both programs so that they utilize pointer notation. (Note that the program shown in Example 9.9 includes the assignment of initial values to individual array elements.) Test both programs using the data given in the examples.

- 10.75** Modify the program given in Example 9.14 (piglatin generator) so that it uses character-type arrays. Modify the program so that it uses pointer notation. Test the program using several lines of text of your own choosing.

- 10.76** Write a complete C program, using pointer notation in place of arrays, for each of the following problems taken from the end of Chap. 9.

- (a) Problem 9.39 (read a line of text, store it within the computer's memory, and then display it backwards).
- (b) Problem 9.40 (process a set of student exam scores). Test the program using the data given in Prob. 9.40.
- (c) Problem 9.42 (process a set of weighted student exam scores, and calculate the deviation of each student's average about the overall class average). Test the program using the data given in Prob. 9.40.
- (d) Problem 9.44 (generate a table of compound interest factors).
- (e) Problem 9.45 (convert from one foreign currency to another).
- (f) Problem 9.46 (determine the capital for a specified country, or the country whose capital is specified). Test the program using the list of countries and their capitals given in Prob. 9.46.
- (g) Problem 9.47(a) (matrix/vector multiplication). Test the program using the data given in Prob. 9.47(a).
- (h) Problem 9.47(b) (matrix multiplication). Test the program using the data given in Prob. 9.47(b).
- (i) Problem 9.47(d) (Lagrange interpolation). Test the program using the data given in Prob. 9.47(d).
- (j) Problem 9.48(a) (blackjack).

- (k) Problem 9.48(b) (roulette).
- (l) Problem 9.48(c) (BINGO).
- (m) Problem 9.49 (encode and decode a line of text).

10.77 Write a complete C program, using pointer notation, that will generate a table containing the following three columns:

$$t \quad ae^{bt} \sin ct \quad ae^{bt} \cos ct$$

Structure the program in the following manner: write two special functions, `f1` and `f2`, where `f1` evaluates the quantity $ae^{bt} \sin ct$ and `f2` evaluates $ae^{bt} \cos ct$. Have `main` enter the values of a , b and c , and then call a function, `table_gen`, which will generate the actual table. Pass `f1` and `f2` to `table_gen` as arguments.

Test the program using the values $a = 2$, $b = -0.1$, $c = 0.5$ where the values of t are $1, 2, 3, \dots, 60$.